

The Krishnaswamy Laboratory
Yale Genetics and Yale SEAS present

Machine Learning for Single Cell Analysis

Online - May 20-29, 2020

When poll is active, respond at **PollEv.com/yaleml**

Text **YALEML** to **22333** once to join

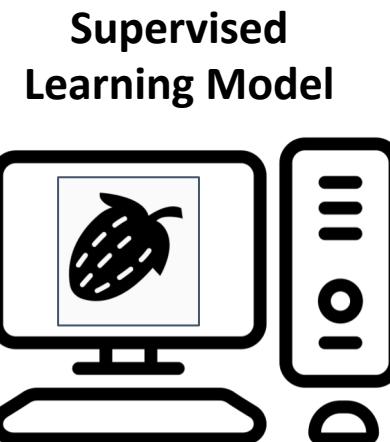
What's the last great TV show or movie you watched?

Recap from Day 1

Two kinds of machine learning

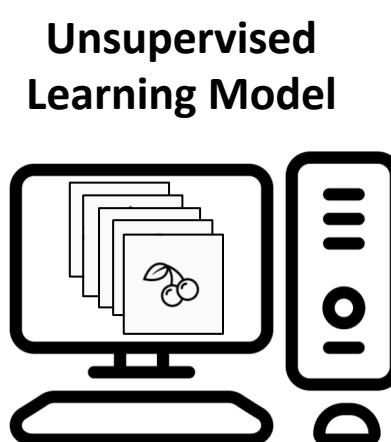
Supervised learning

- Have a bunch of labelled data, want to label new data
- Learn a function $f(X) \rightarrow Y$ where all values of Y are known for some samples of X



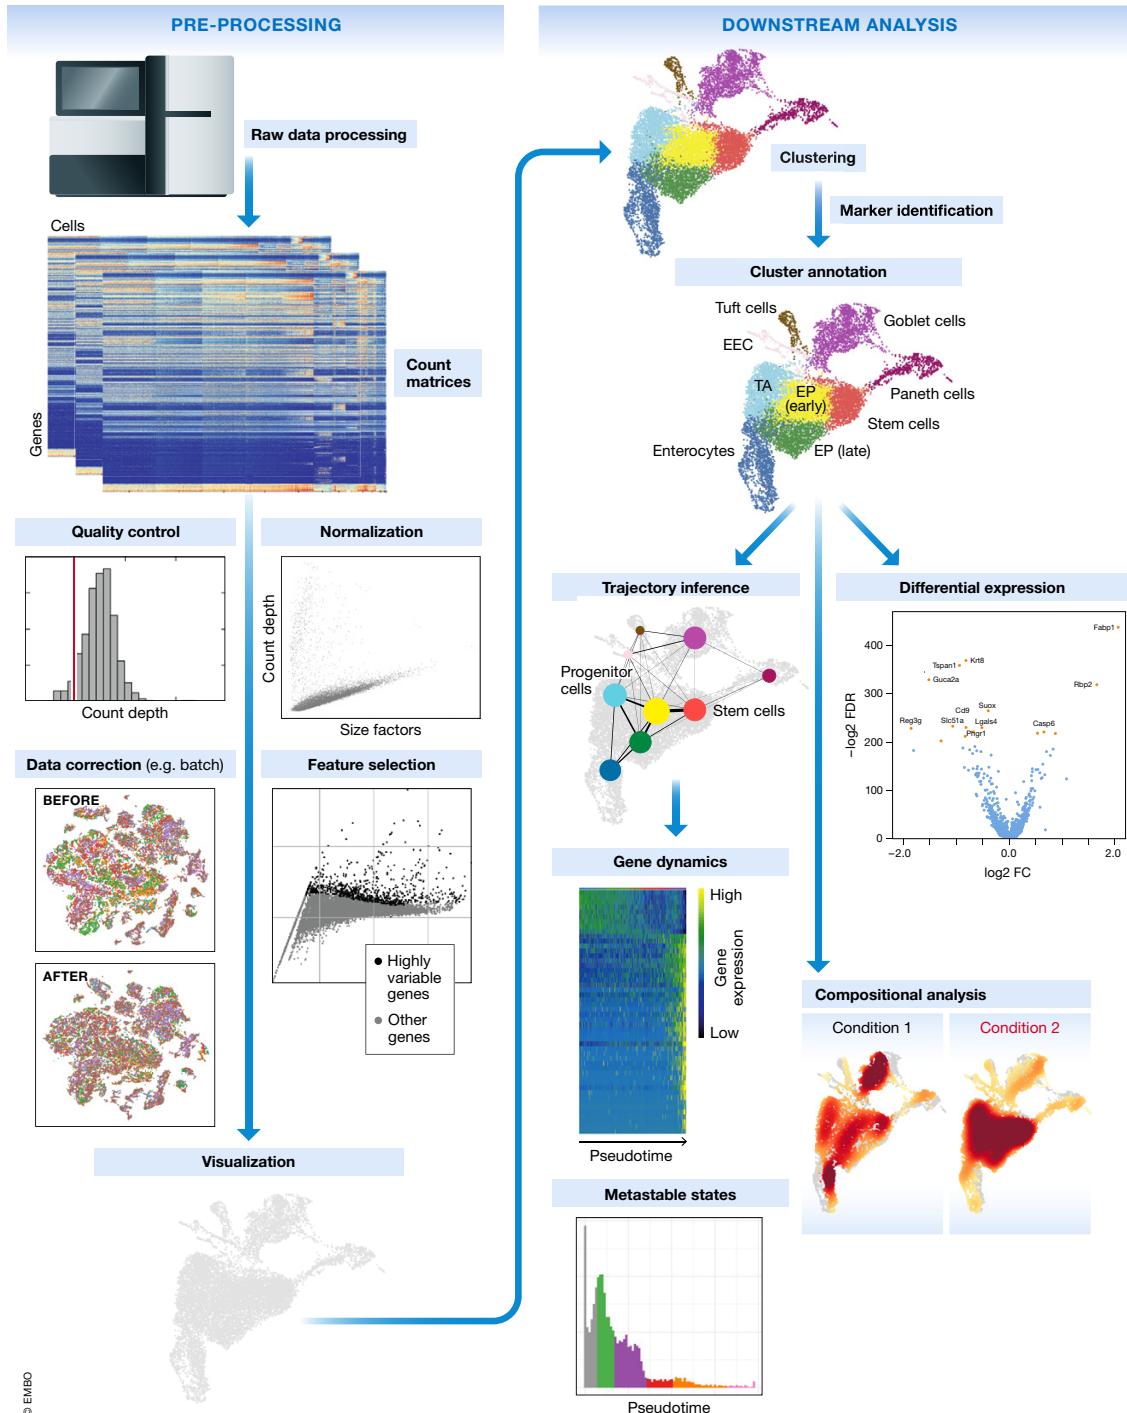
Unsupervised learning

- Have a bunch of unlabeled data, want to organize it
- Learn an embedding $f(X) \rightarrow Y, X \in \mathbb{R}^n, Y \in \mathbb{R}^m, n \gg m$
- Lower dimensional, easier to interpret (e.g. as clusters)

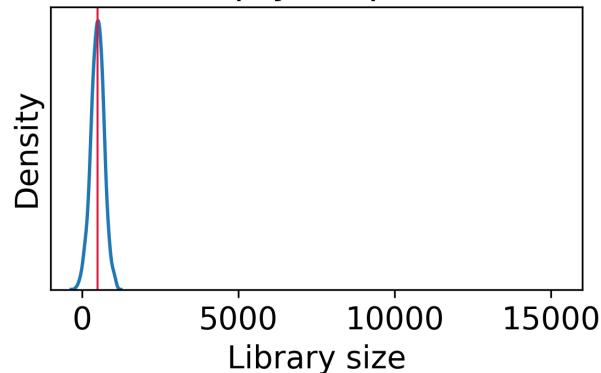


Standard Single-Cell RNA-seq Workflow

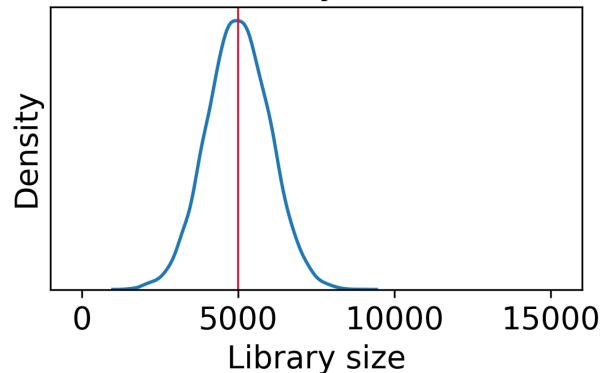
1. Sequencing and read mapping
2. Quality control and filtering
3. Normalization
4. Data Correction
5. Dimensionality reduction and visualization
6. Downstream analysis
 1. Clustering
 2. Trajectory inference
 3. Differential expression
7. Comparison of multiple conditions



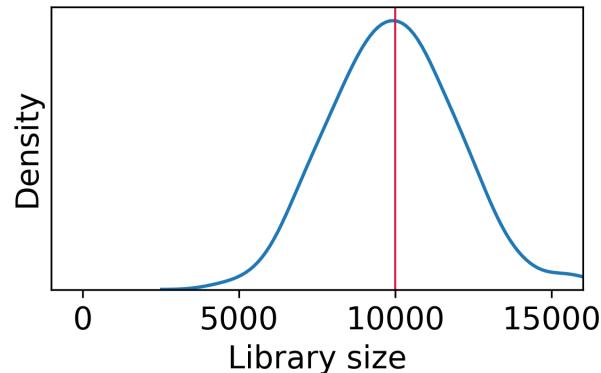
Empty droplets



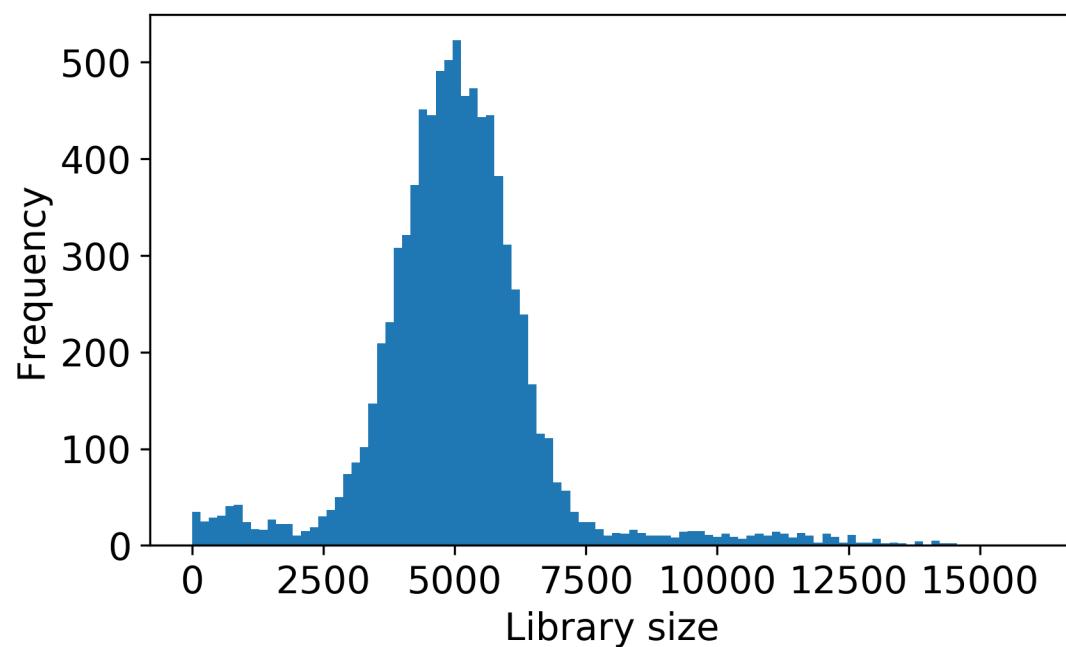
Healthy cells



Doublets



9K simulated cells



Course Schedule		
Day 1 – Wednesday, May 20th		
Lecture	Download from GitHub	Introduction to scRNA-seq and Preprocessing
Exercise	Run in Google Colab	1.0. Preprocessing Embryoid Body Data (Beginner)
	Run in Google Colab	1.0. Preprocessing Embryoid Body Data (Advanced)
	Run in Google Colab	1.0. Preprocessing Embryoid Body Data (Answer Key) 
	Run in Google Colab	1.1. Loading and pre-processing your own data (optional)
Day 2 – Thursday, May 21st		
Lecture	View on Google Drive	Manifold Learning and Dimensionality Reduction
Exercise	Run in Google Colab	2.0. Plotting UCI Wine Data
	Run in Google Colab	2.1. Learning Graphs from Data
	Run in Google Colab	2.2. Visualizing UCI Wine Data
	Run in Google Colab	2.3. PCA on Retinal Bipolar Data
	Run in Google Colab	2.4. Visualizing Retinal Bipolar Data
	Run in Google Colab	2.5. Visualizing Embryoid Body Data (Advanced)
Day 3 – Friday, May 22nd		
Lecture	View on Google Drive	Clustering and Data Denoising
	Run in Google Colab	

<https://www.krishnaswamylab.org/workshop>

KL

Krishnaswamy Lab Help ▾

• Daniel Burkhardt

Y2

Threads

Mentions & reactions

Drafts

Saved items

Channel browser

People

Apps

Files

Show less

Channels

2020-workshop-byod-help

2020-workshop-coding-help

2020-workshop-group-01

2020-workshop-group-08

2020-workshop-group-09

2020-workshop-group-15

2020-workshop-group-25

2020-workshop-main

2020-workshop-math-help

Search Krishnaswamy Lab Help

#2020-workshop-main ☆

8 118 | 1 | <https://zoom.us/j/92462966236?pwd=d3U1Y0hpQ2JYcFFOZWxmQlI4WnpEdz09>

Details

Friday, May 15th

Monday, May 18th

Pinned by you

Daniel Burkhardt 11:53 AM

Hi everyone! Welcome to the main channel for the 2020 Machine Learning for Single Cell Analysis Workshop! Please join the following channels:

1. #2020-workshop-coding-help
2. #2020-workshop-math-help
3. #2020-workshop-byod-help
4. #2020-workshop-misc-help

2 1

Yesterday

Scott Gigante 12:19 PM

To respond to the poll, please go to <https://PollEv.com/yaleml>

PollEv.com
Poll Everywhere - Audience Participation Site

This is the place to be if you're trying to participate in a live poll.

1 1

Message #2020-workshop-main

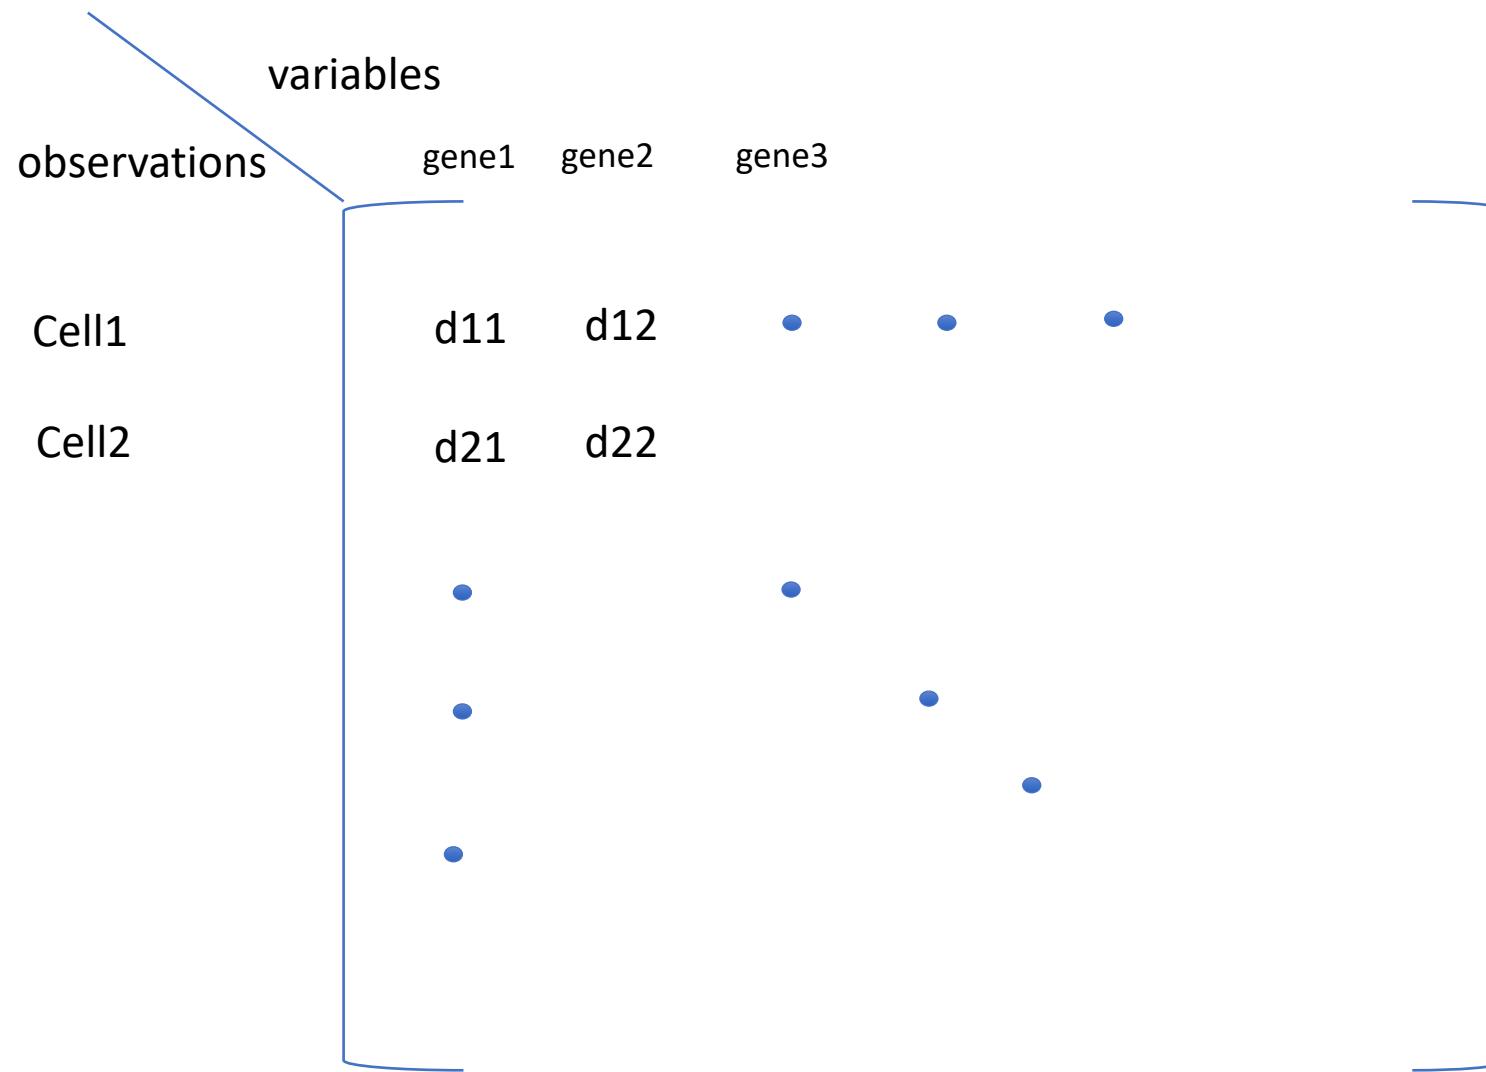
Text input field with rich text toolbar

Aa @ ☺ ⌂ ➤

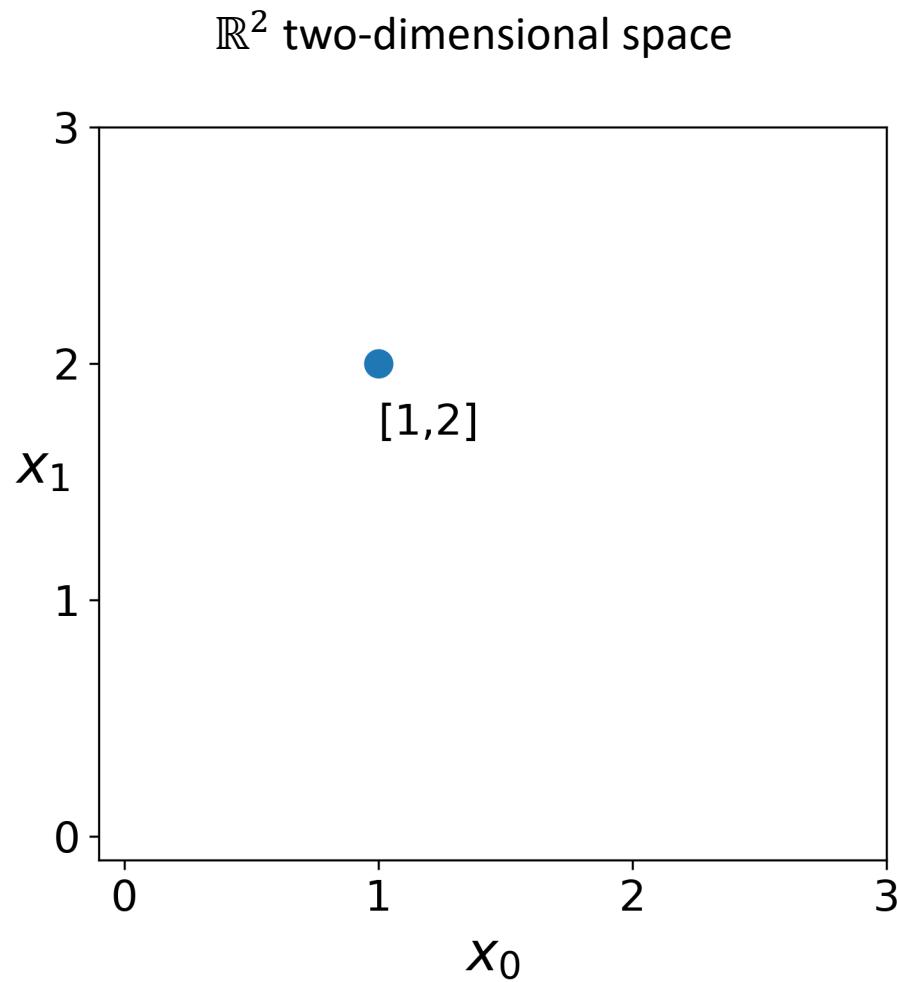
Day 2: Visualization, Dimensionality Reduction and Manifold Learning

Thinking about high dimensional data

Our Data is a High-dimensional Matrix



What does “dimensionality” mean?



Features are dimensions

vector coordinate space

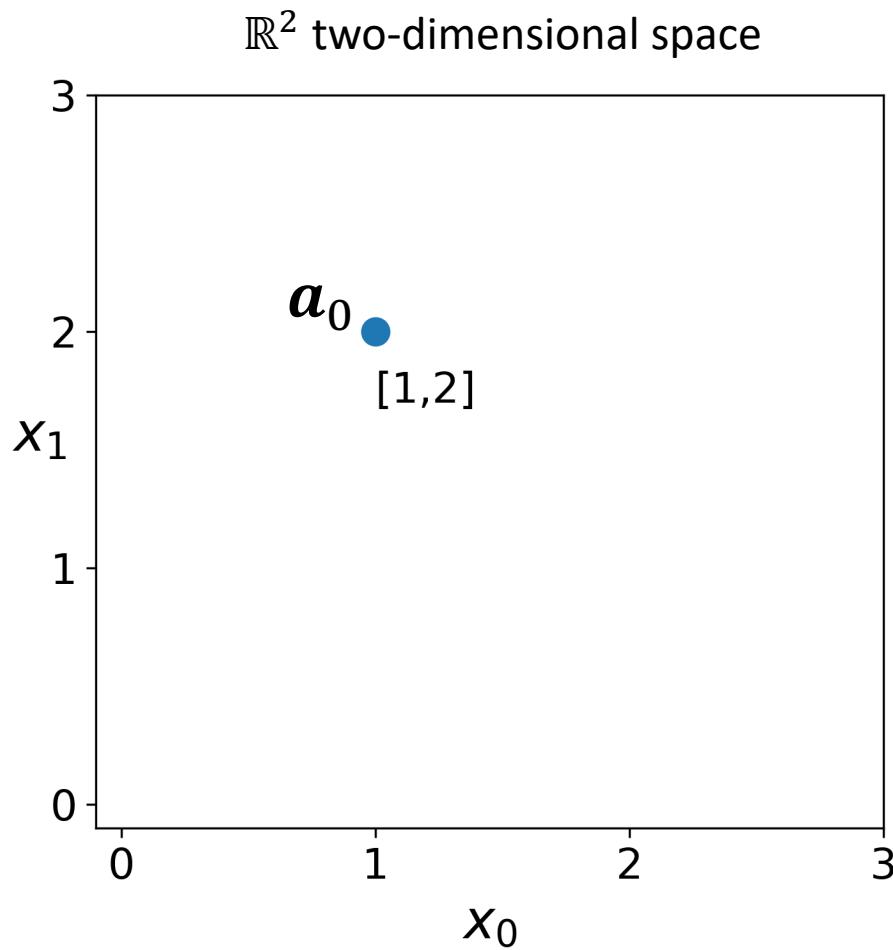
$a_0 \in \mathbb{R}^2$

$a_0 = [1, 2]$

$x_0 \quad x_1$

features

$a_{[:,0]}$



Features are dimensions

$$\mathbf{a} \in \mathbb{R}^2$$

$$\mathbf{a} = [1, 2]$$

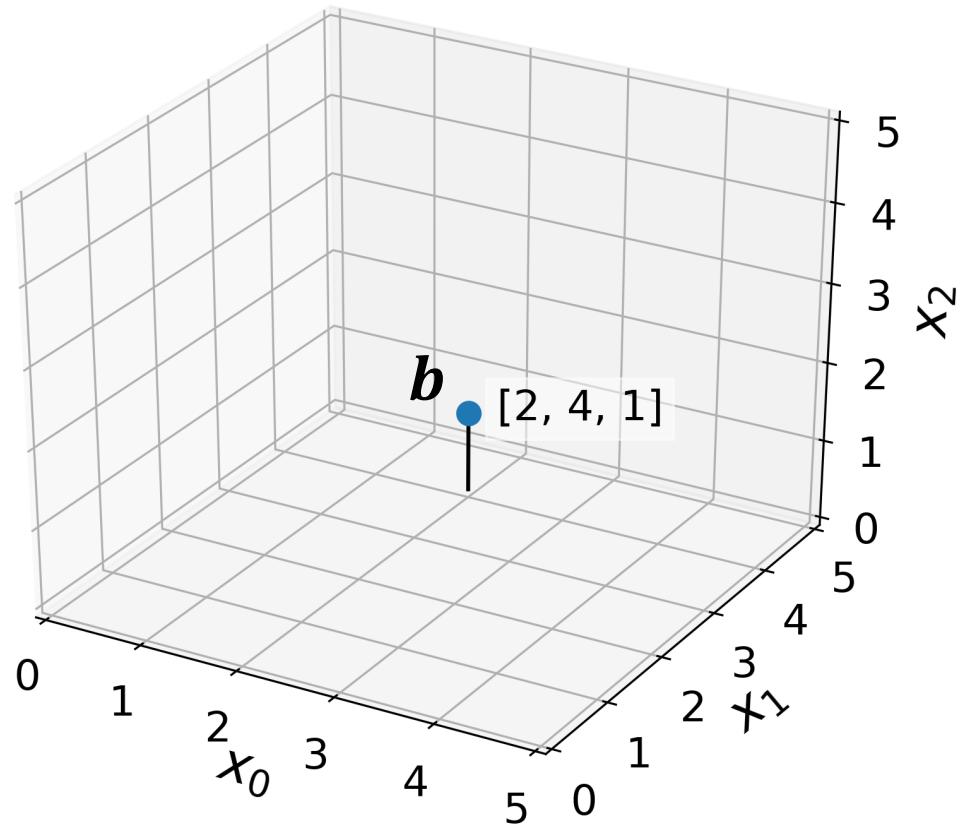
x_0 x_1

$$\mathbf{b} \in \mathbb{R}^3$$

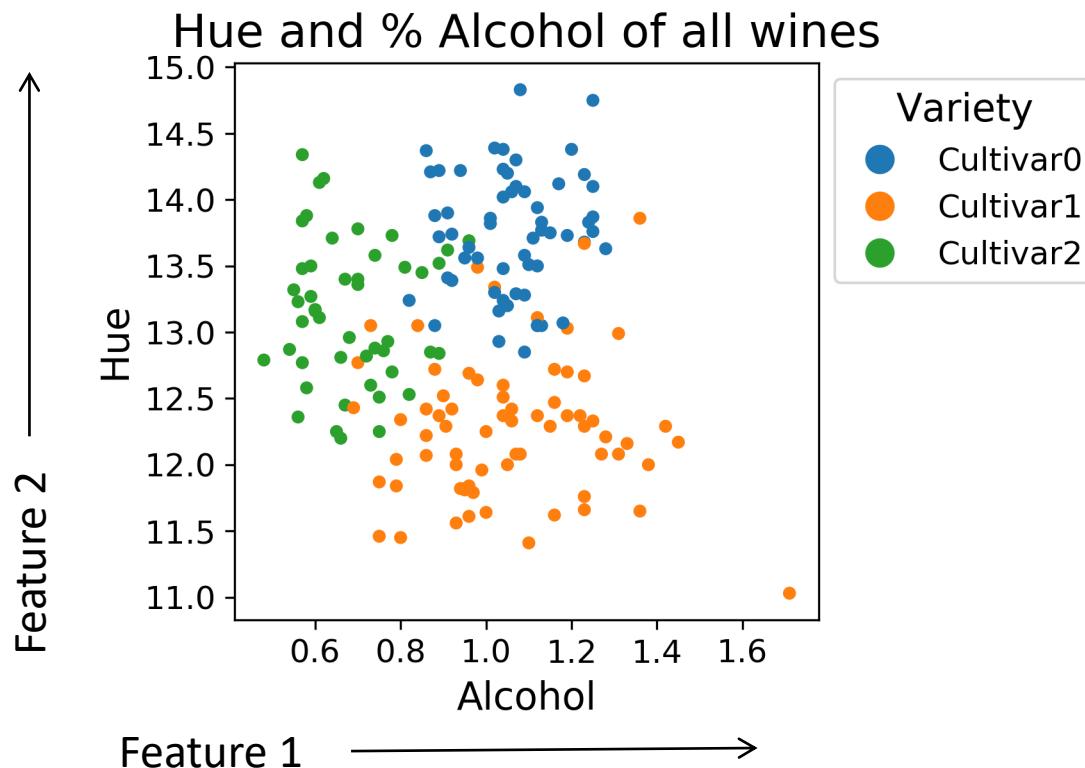
$$\mathbf{b} = [2, 4, 1]$$

x_0 x_1 x_2

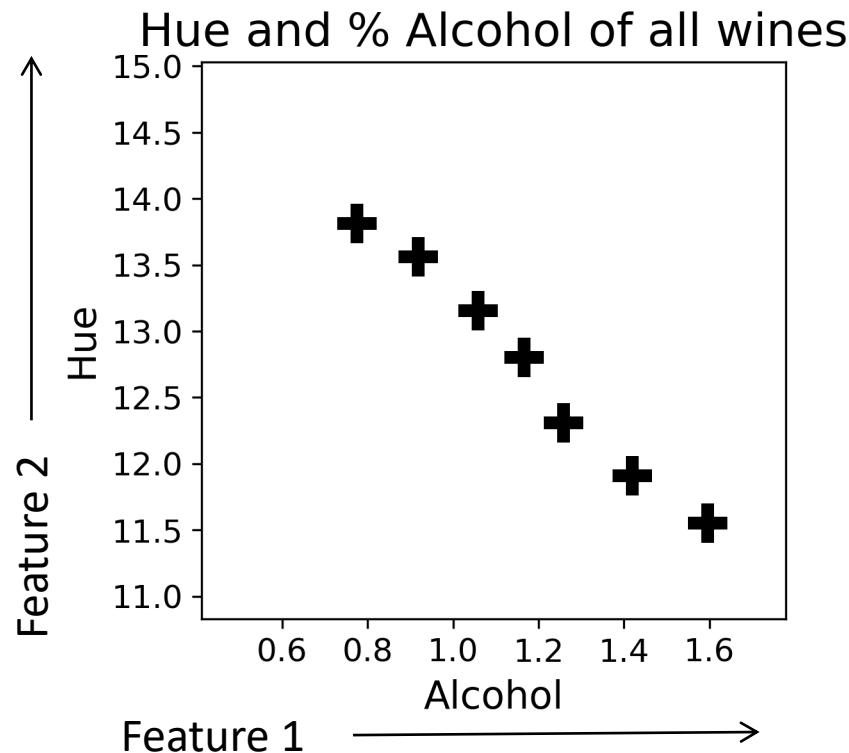
\mathbb{R}^3 three-dimensional space



Features reveal structure

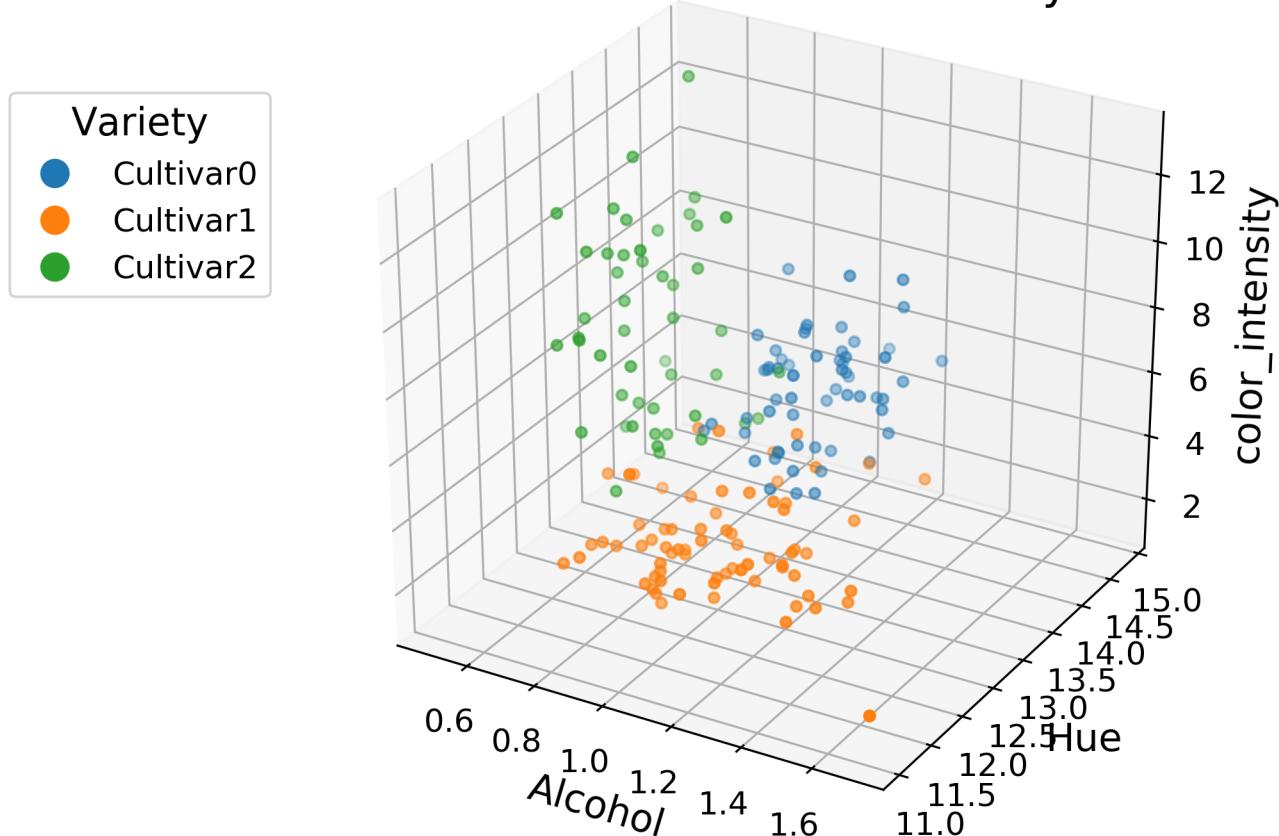


Many types of structure can exist in data

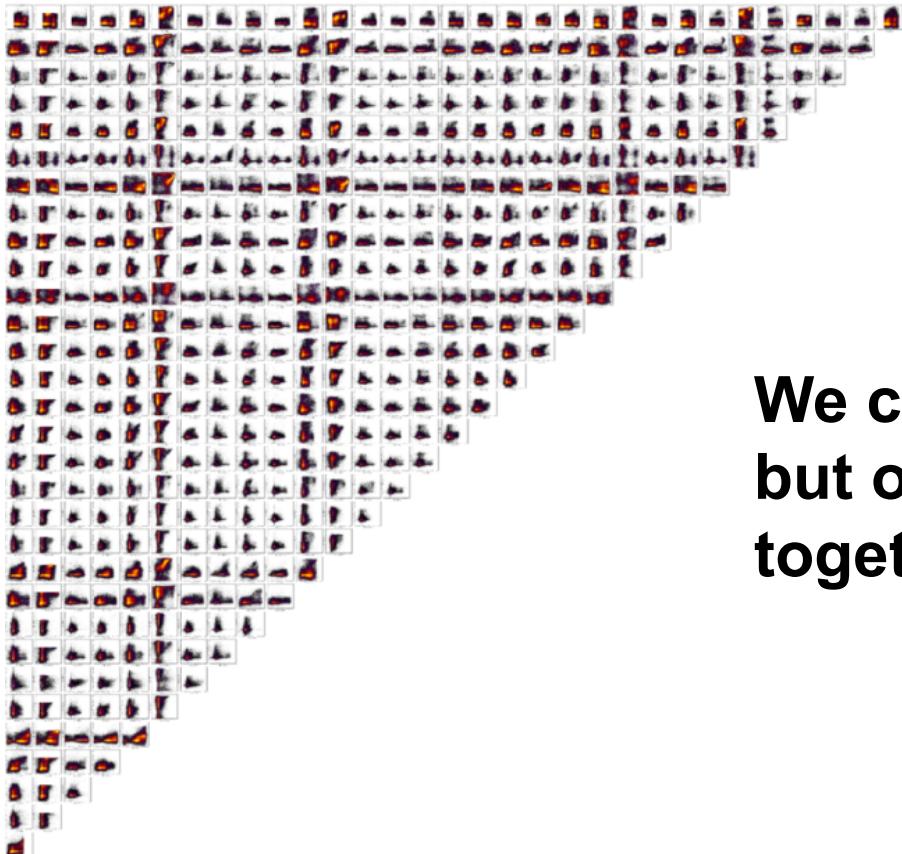


More Dimensions = More Accurate Structure

Hue and % Alcohol and Color Intensity of all wines



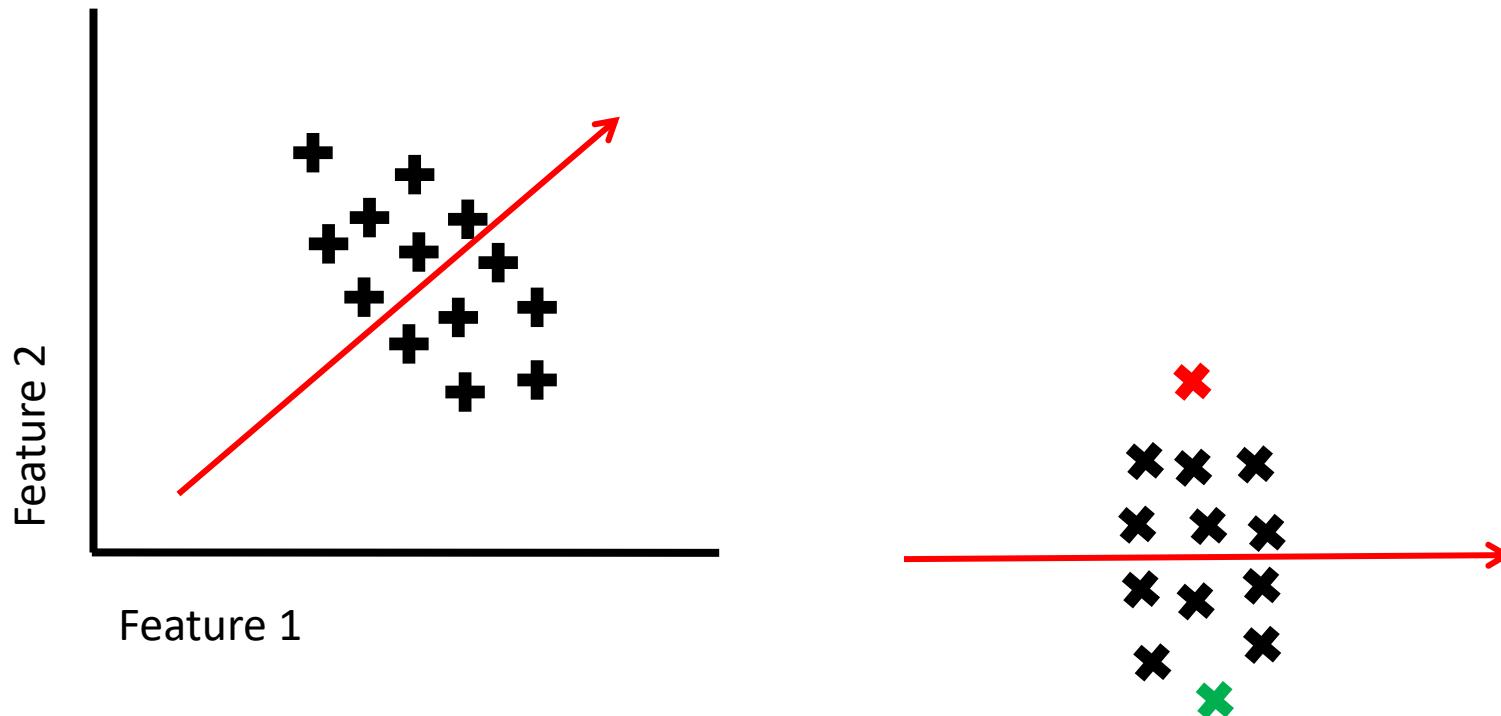
Problem: We can only see in 3D (not 20K D)



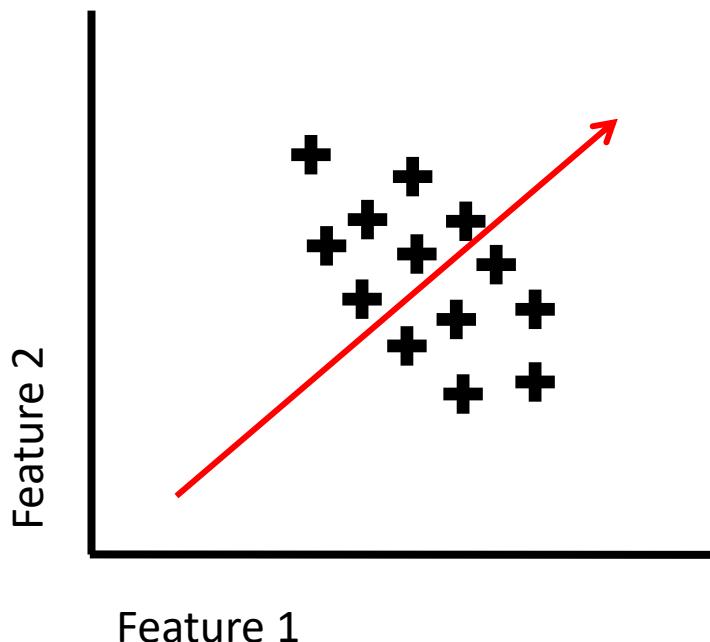
**We could try all 2D pairs,
but our brains can't put this
together**

Solution: Dimensionality Reduction

Thought Exercise: Reduce this to 1D



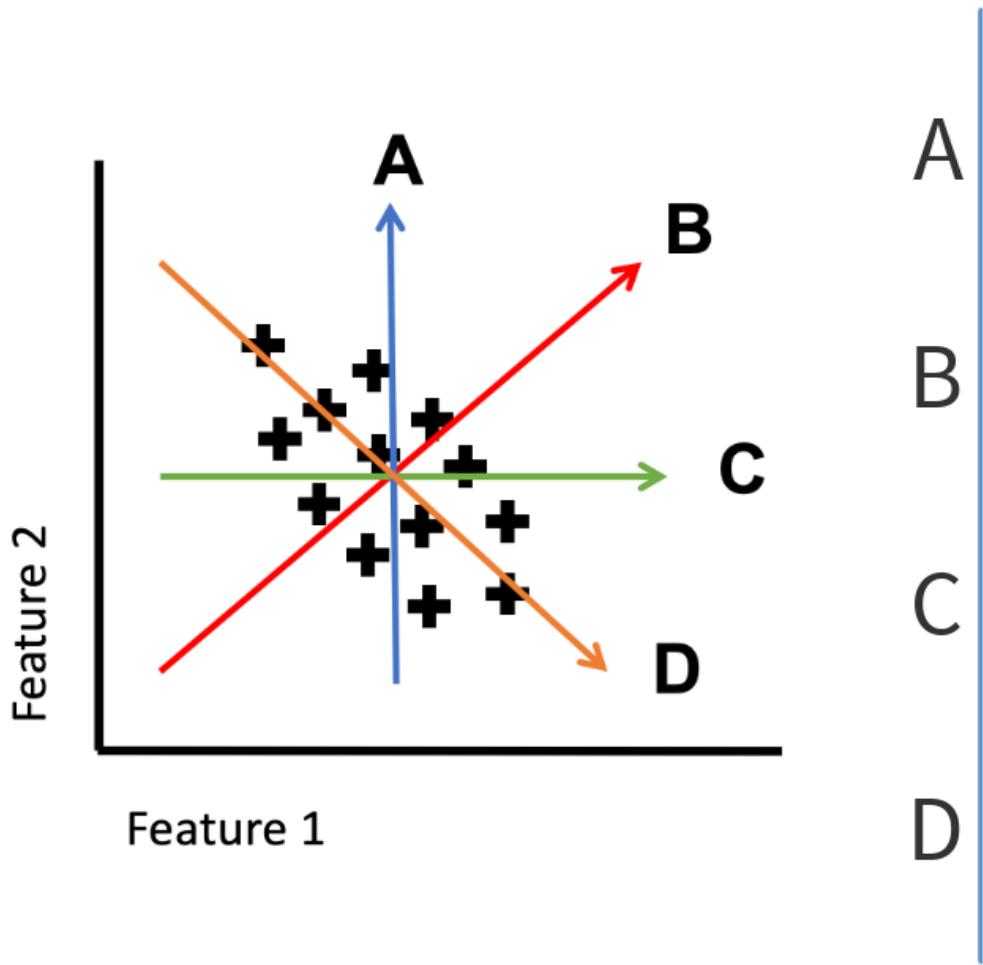
Thought Exercise: Reduce this to 1D



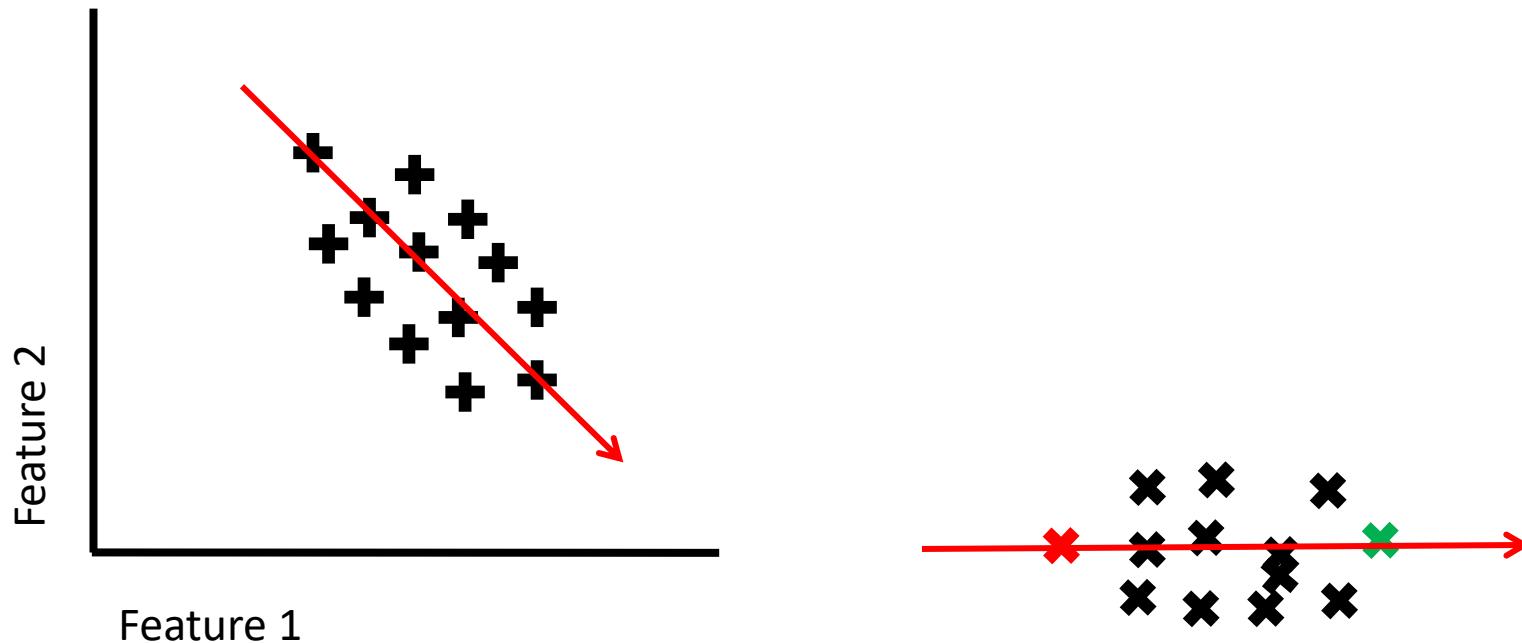
Most distant points on top of each other !



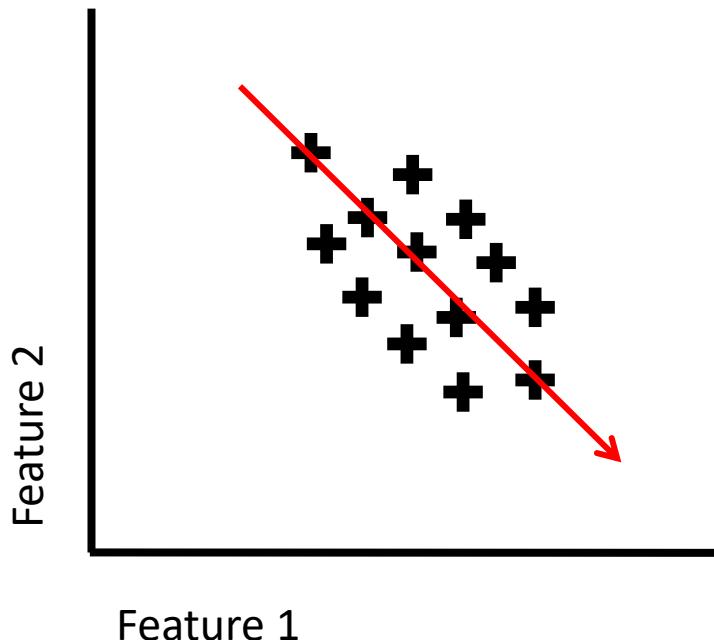
Which line do you think would be the best line to consider the data onto?



Thought Exercise: Reduce this to 1D



Thought Exercise: Reduce this to 1D

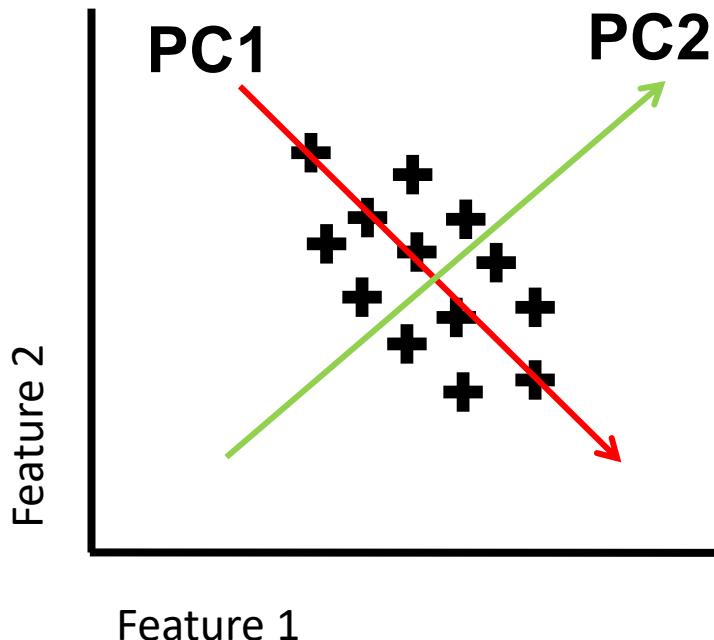


What is special about this line?

It explains most of the variance in the data



Principle Components Analysis



PCA linear directions
in the data that explain
the most variance

PC1 explains the most

PC2 explains the next most
and is orthogonal to PC1

How do we find these directions?

Covariance Matrix

$$\mathbf{X} = (X_1, X_2, \dots, X_n)^T$$

$$K_{X_i X_j} = \text{cov}[X_i, X_j] = E[(X_i - E[X_i])(X_j - E[X_j])]$$

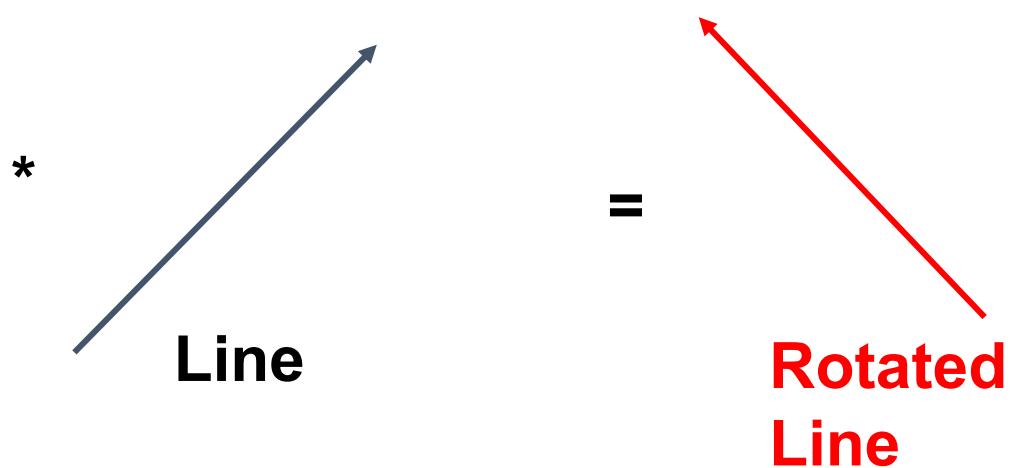
$$K_{\mathbf{X}\mathbf{X}} = \begin{bmatrix} E[(X_1 - E[X_1])(X_1 - E[X_1])] & E[(X_1 - E[X_1])(X_2 - E[X_2])] & \cdots & E[(X_1 - E[X_1])(X_n - E[X_n])] \\ E[(X_2 - E[X_2])(X_1 - E[X_1])] & E[(X_2 - E[X_2])(X_2 - E[X_2])] & \cdots & E[(X_2 - E[X_2])(X_n - E[X_n])] \\ \vdots & \vdots & \ddots & \vdots \\ E[(X_n - E[X_n])(X_1 - E[X_1])] & E[(X_n - E[X_n])(X_2 - E[X_2])] & \cdots & E[(X_n - E[X_n])(X_n - E[X_n])] \end{bmatrix}$$

Matrices as Transformations

- Matrices only store data, but can also represent ***transformations***
- Specifically they are linear transformations
- They transform lines by changing their length and angle from origin

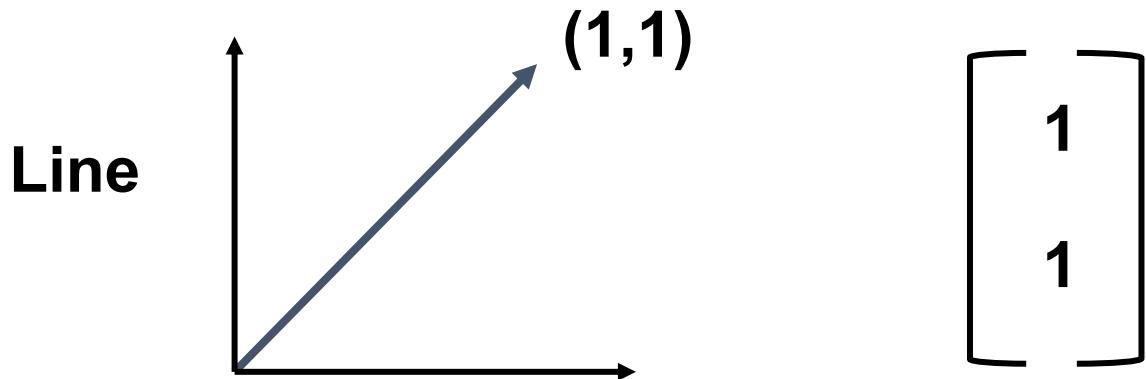
$$\begin{bmatrix} \cos(90) & -\sin(90) \\ \sin(90) & \cos(90) \end{bmatrix}$$

Rotation matrix



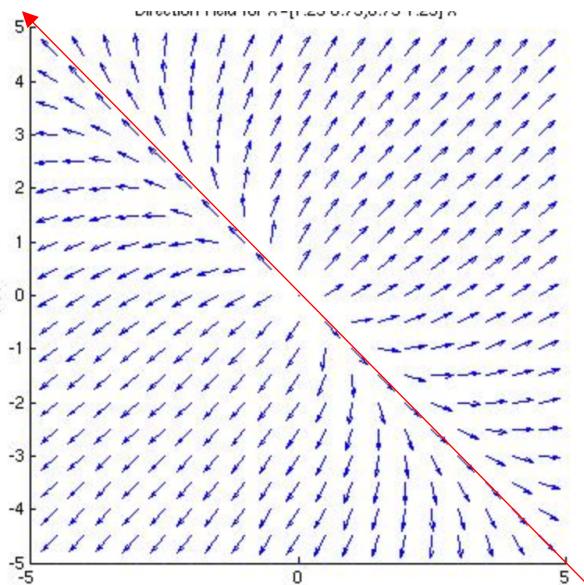
Matrix–Vector Notation

- Lines can be described as vectors from the origin



$$\begin{bmatrix} \cos(90) & -\sin(90) \\ \sin(90) & \cos(90) \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

Eigenvectors



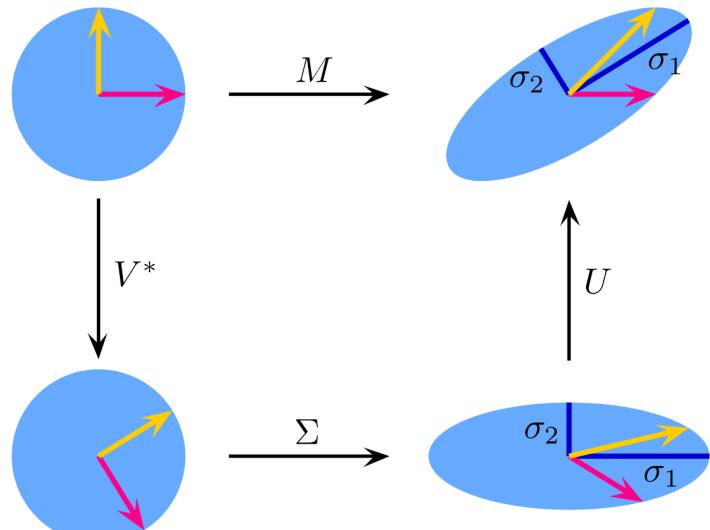
Rotation matrices rotate lines in only one direction

Other matrices can move different lines in different directions

- Such linear transformations have fixed points called *Eigenvectors*
- In other words, the transformation only **stretches** the vector and does not rotate it.

Eigendecomposition

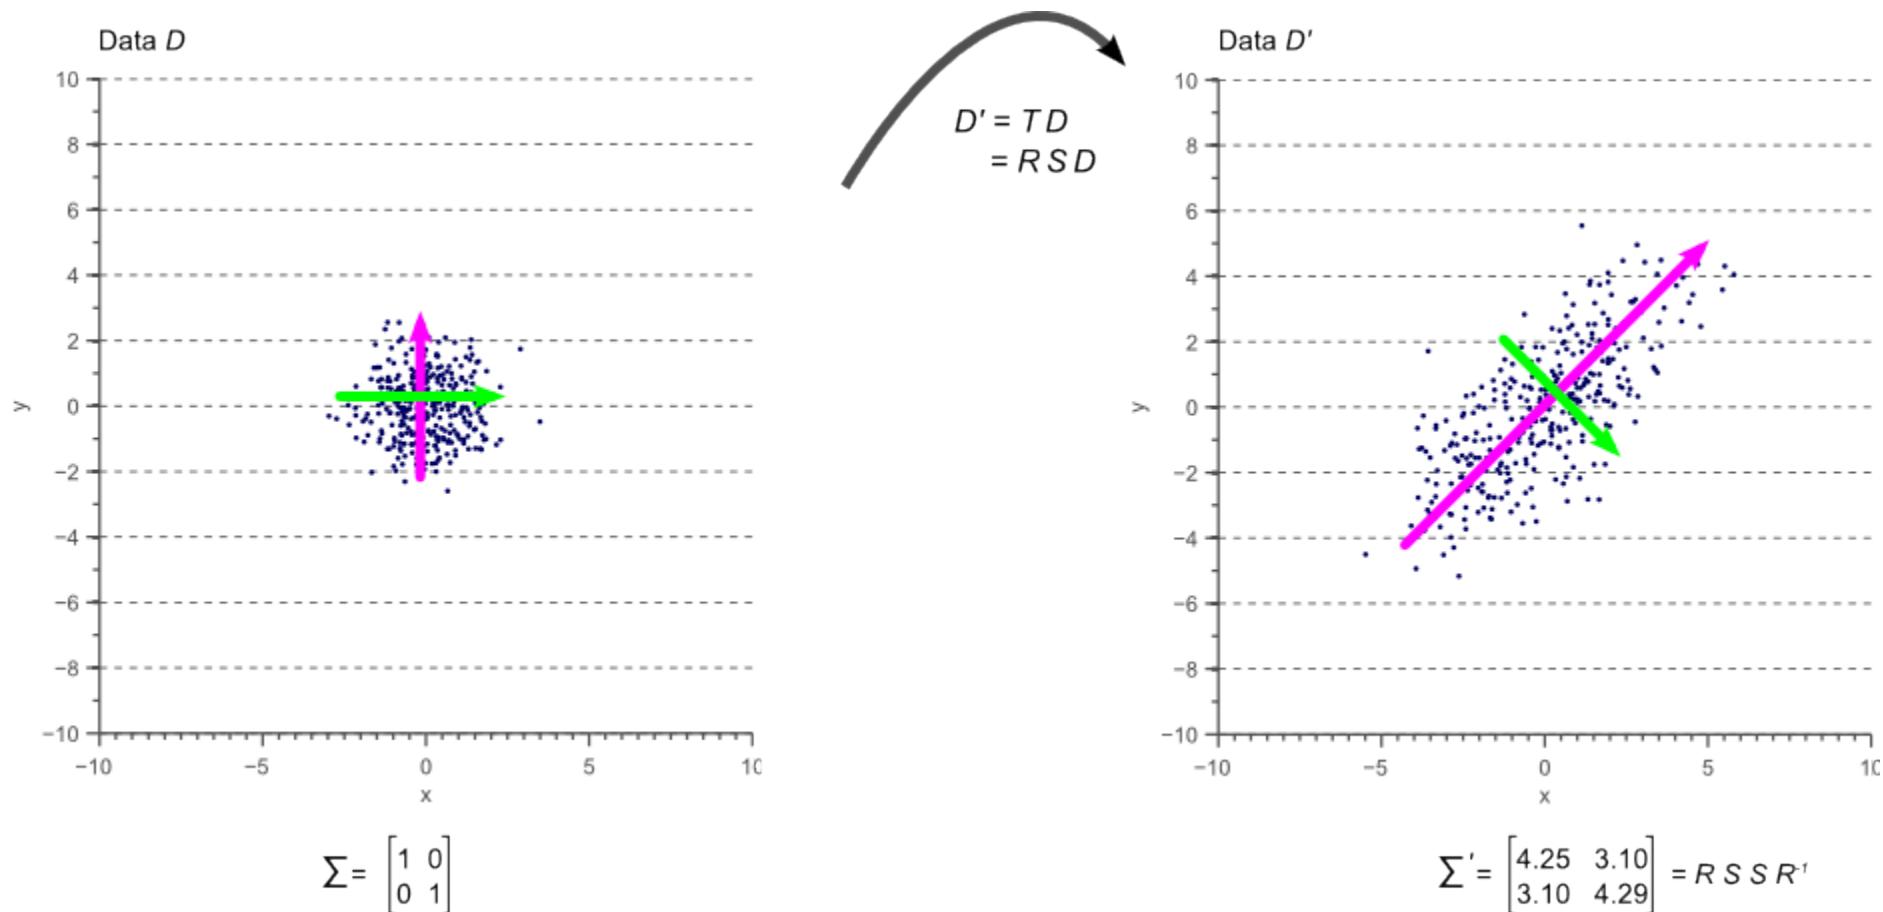
- Eigenvectors can be used to decompose a linear transformation into a rotation + stretch + anti-rotation



Non-square matrices have similar OP called **Singular value decomposition**

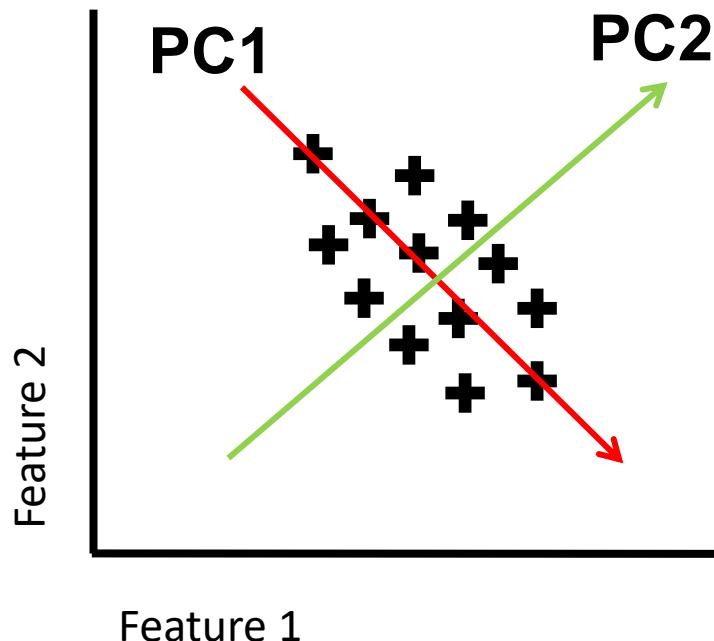
$$M = U \cdot \Sigma \cdot V^*$$

Eigenvectors of Covariance



Matrix creates correlation structure of data on unit blob

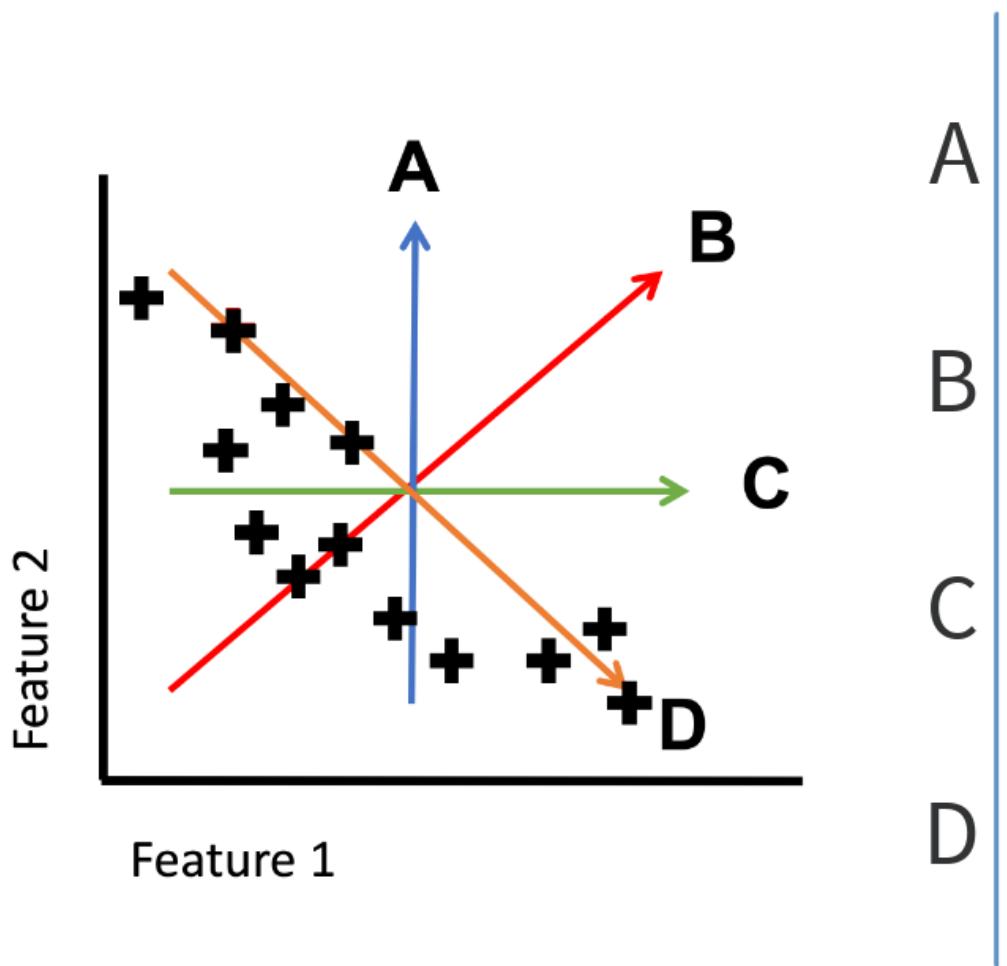
Principle Components Analysis



PC1 is first eigenvector (with
Highest eigenvalue)

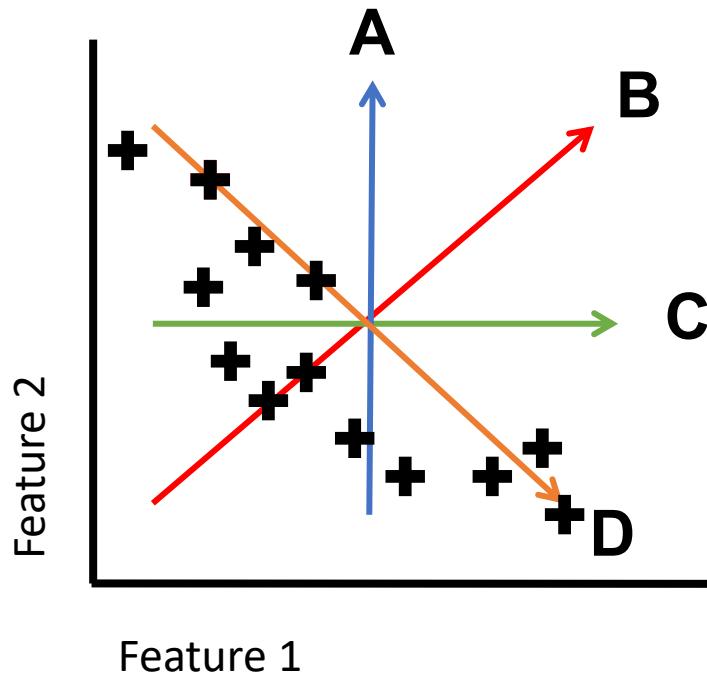
PC2 is next eigenvector

Which line would be PC1 for the following data?

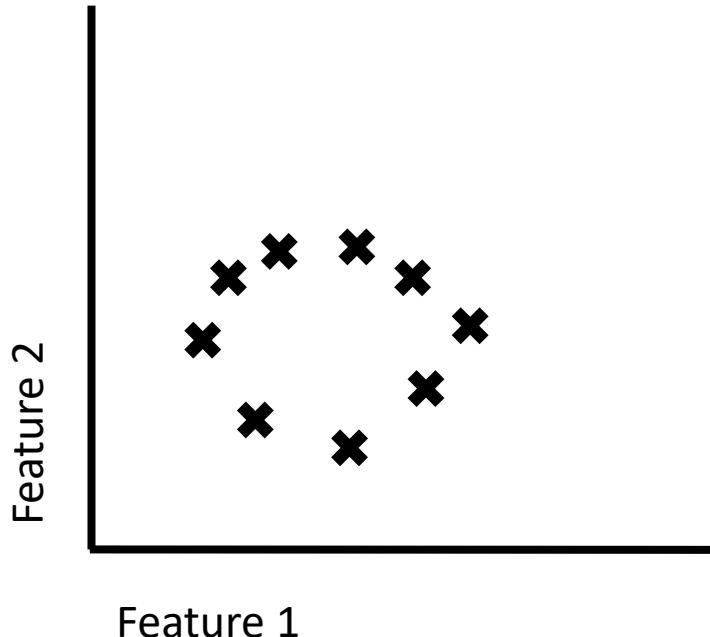


Quick quiz

Which line would be PC1 for the following data?



Non-linear structure



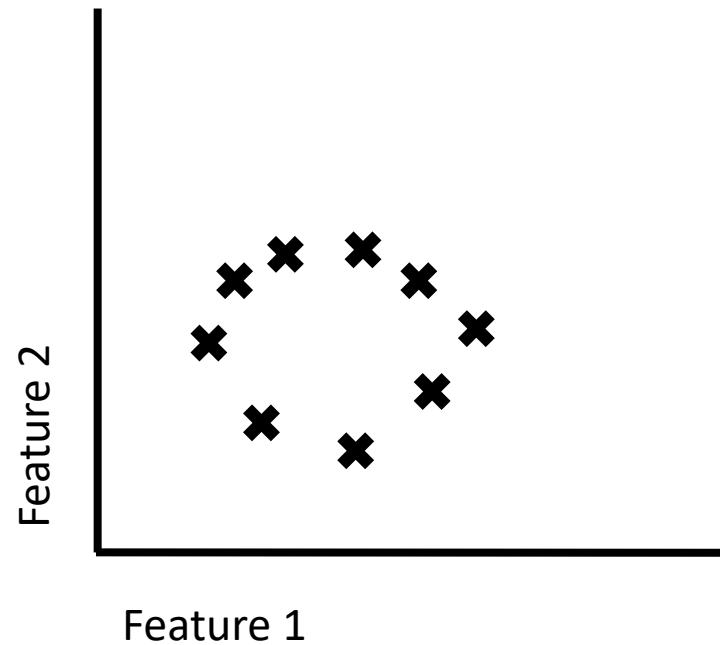
What if the data looked like this?

What line would be good
to project to here?

Non-linear dimensionality reduction

Non-linear dimensions

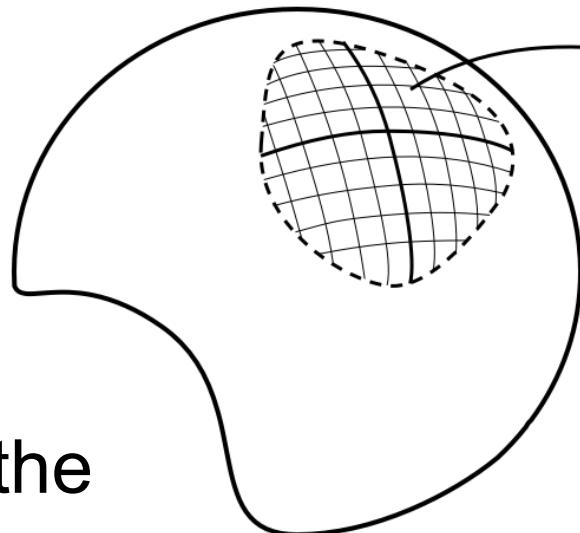
- To find non-linear or “coiled” dimensions in a dataset we have to understand, and describe the **shape** of the data
- The data can live in lower dimensions
- This can be abstractly called the data “manifold”



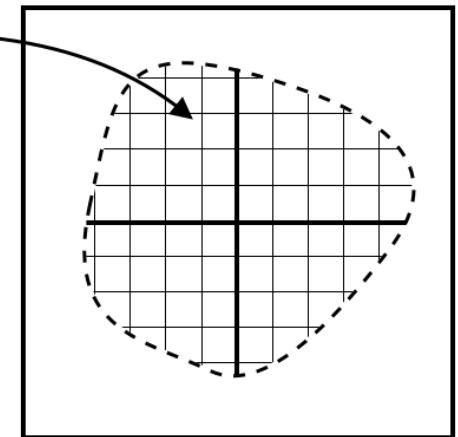
What is a manifold?

- Locally smooth
- Locally Euclidean
- Generally, lower dimensional than the ambient space (i.e. a subspace)

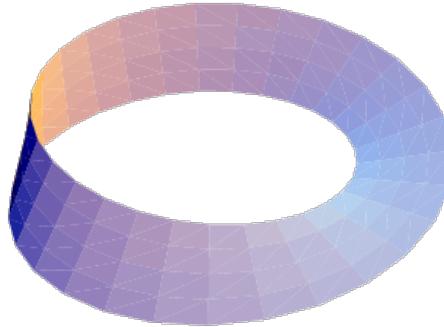
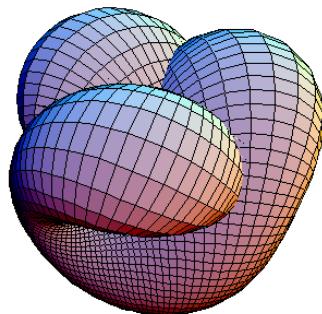
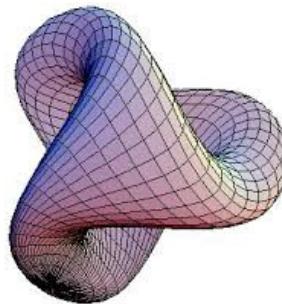
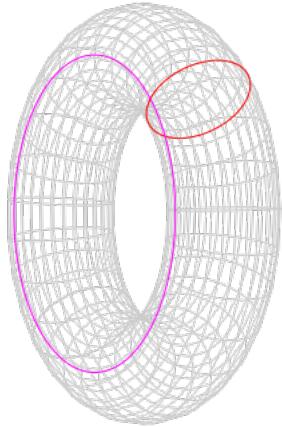
Surface in \mathbb{R}^3



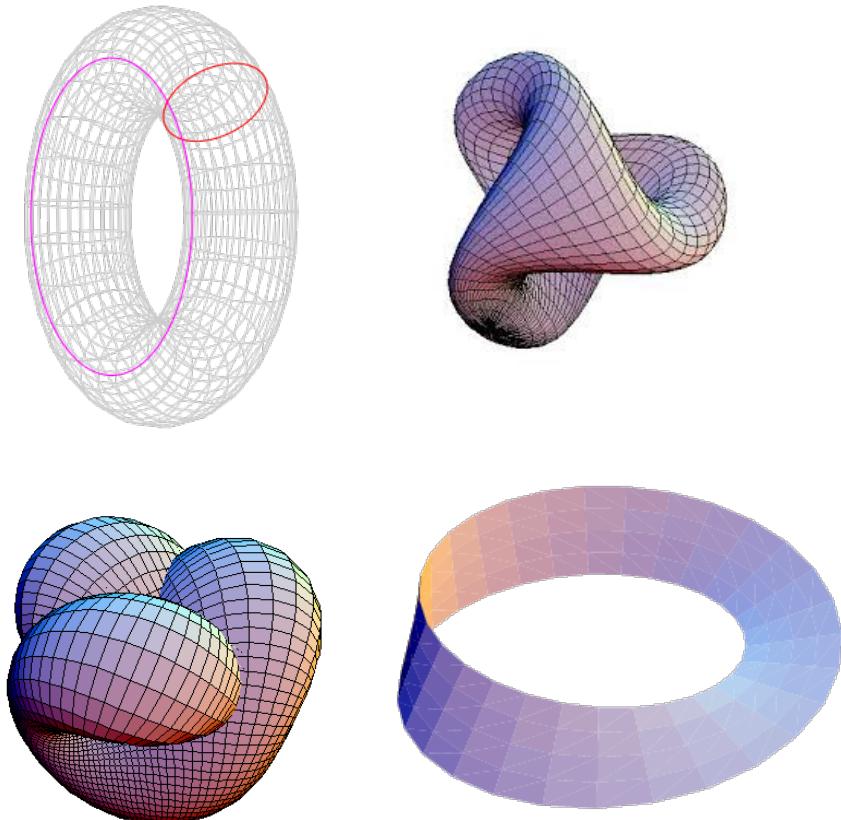
Local view in \mathbb{R}^2



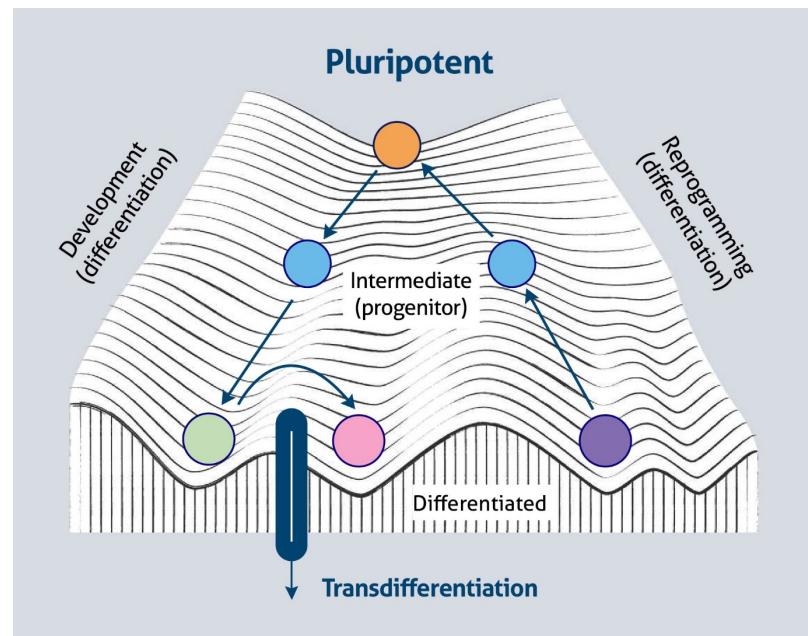
Examples of manifolds



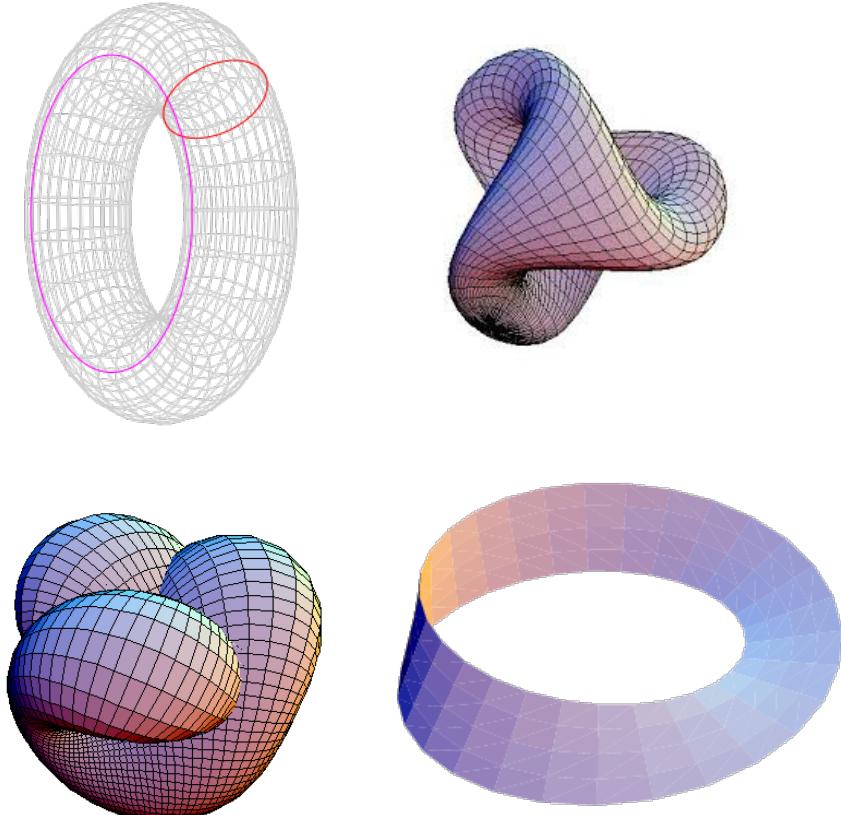
Examples of manifolds



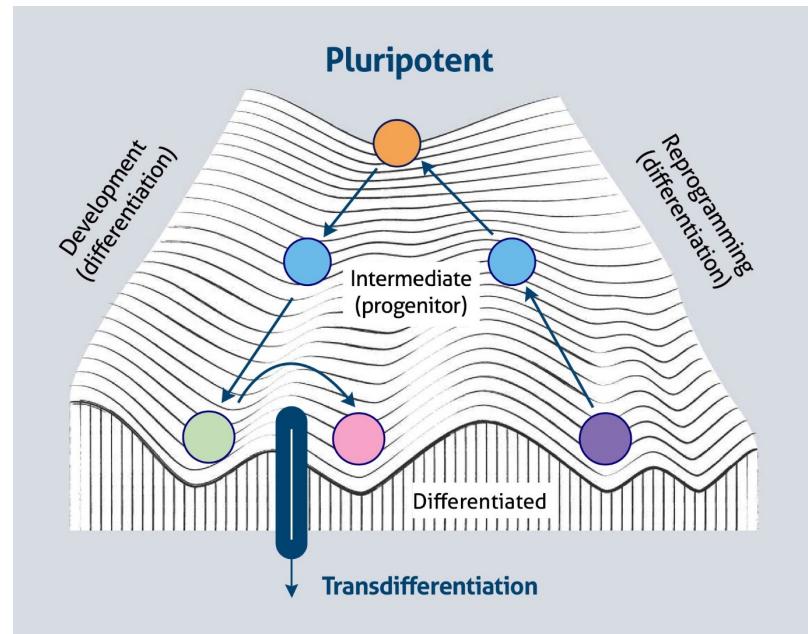
Waddington's landscape



Examples of manifolds



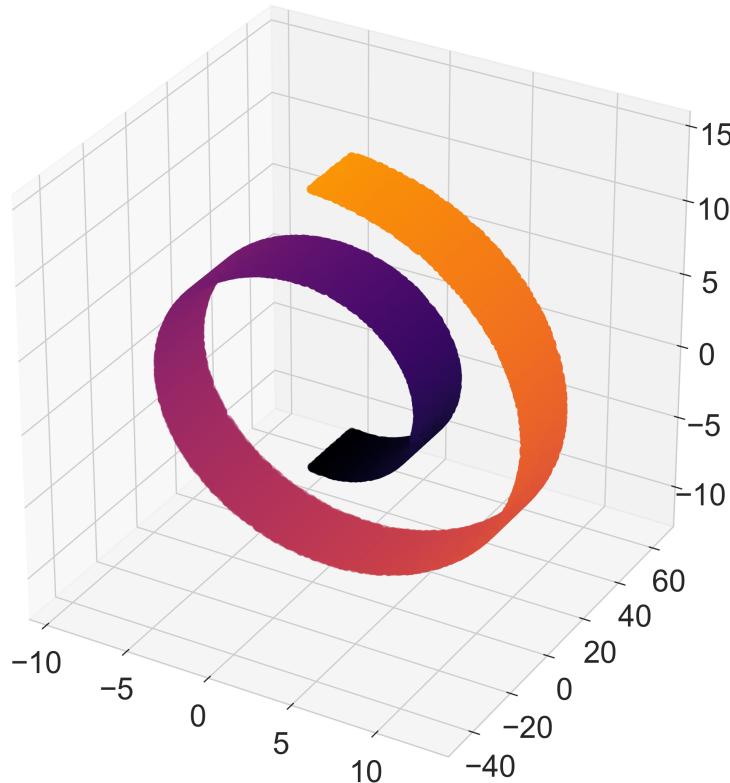
Waddington's landscape



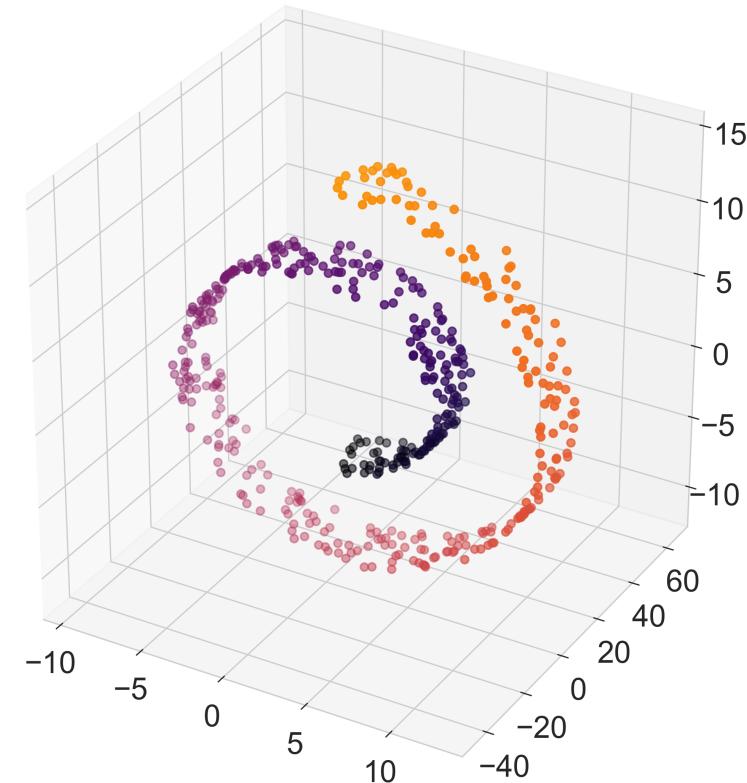
- Smooth transitions between states
- Small changes in gene expression are linear

The swiss roll is a plane (\mathbb{R}^2) embedded in \mathbb{R}^3

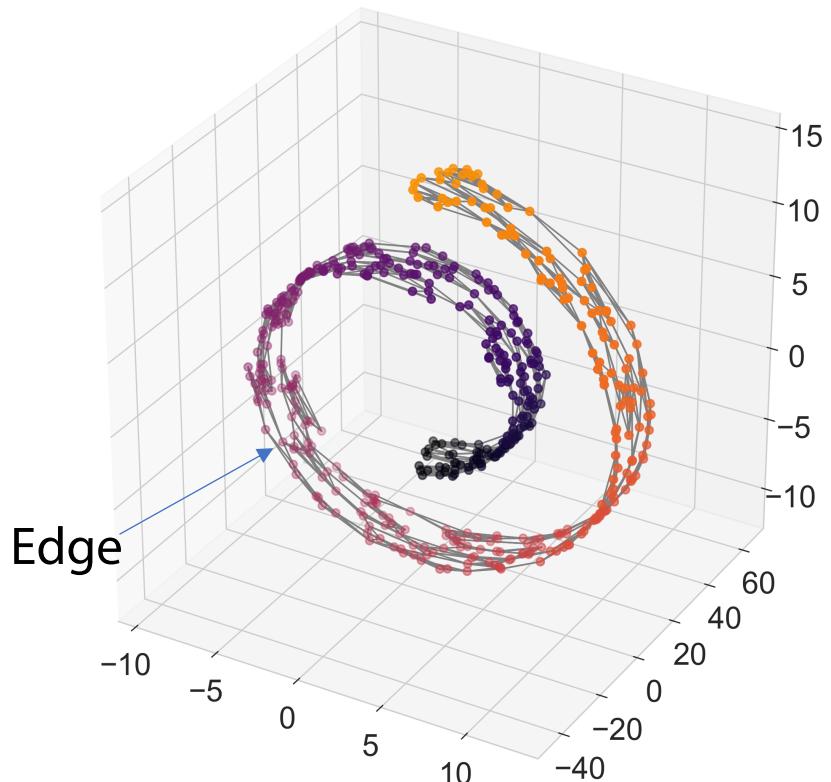
True manifold



Data sampled from manifold



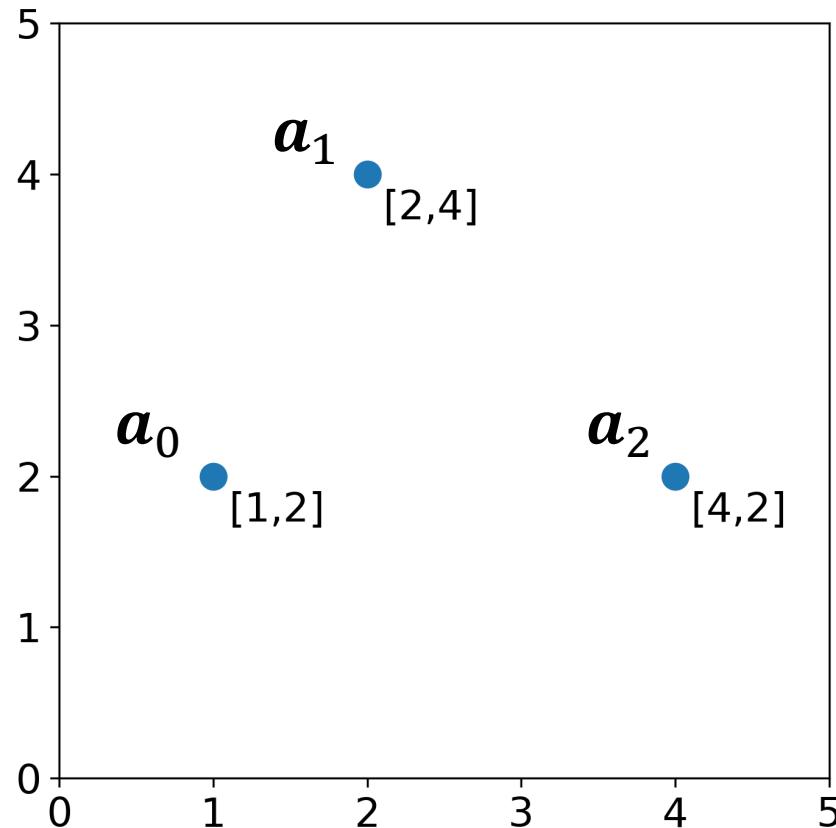
How do we find the data manifold?



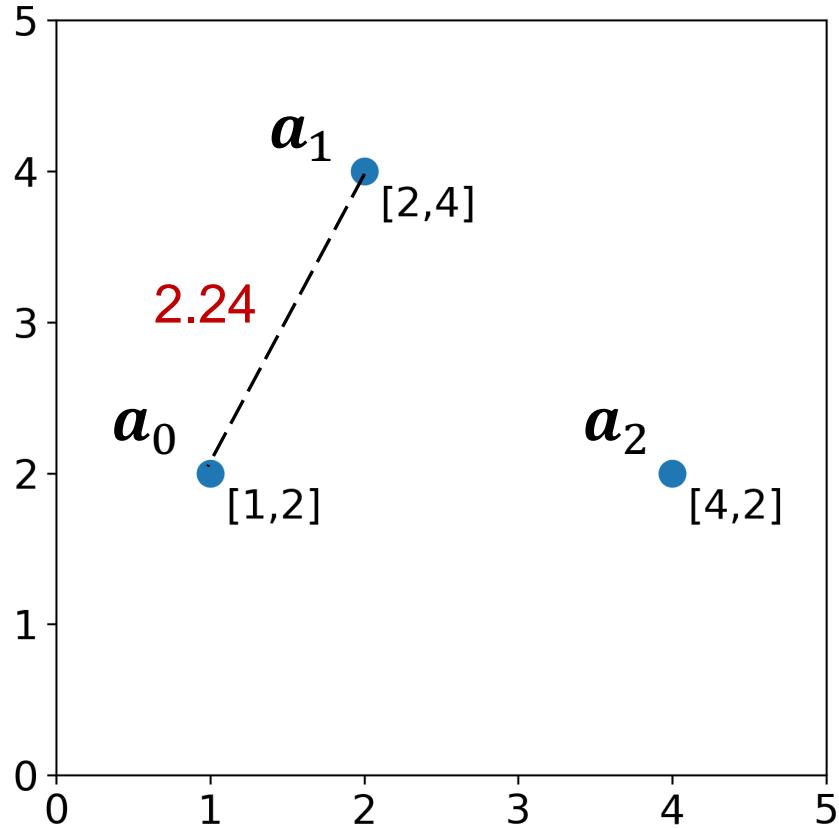
If you connect the nearest neighbors of any data point with a line (also called EDGE) it forms a mesh of the data that can be unrolled to find the shape.

This is called a nearest neighbors graph

How far away is each point from a_0 ?

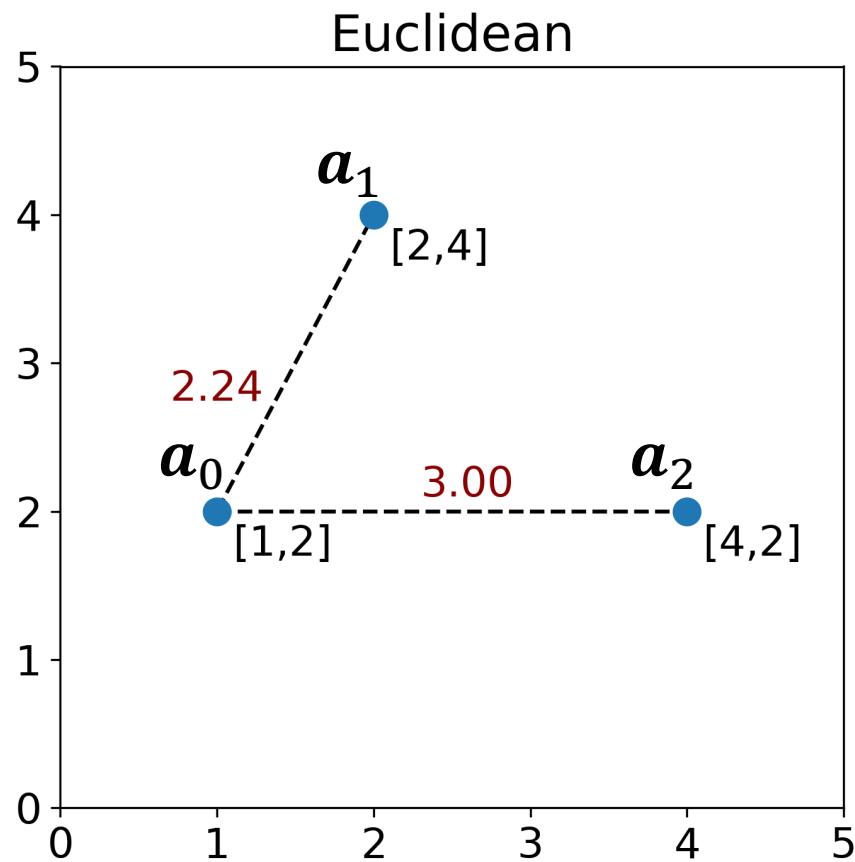


Euclidean Distance

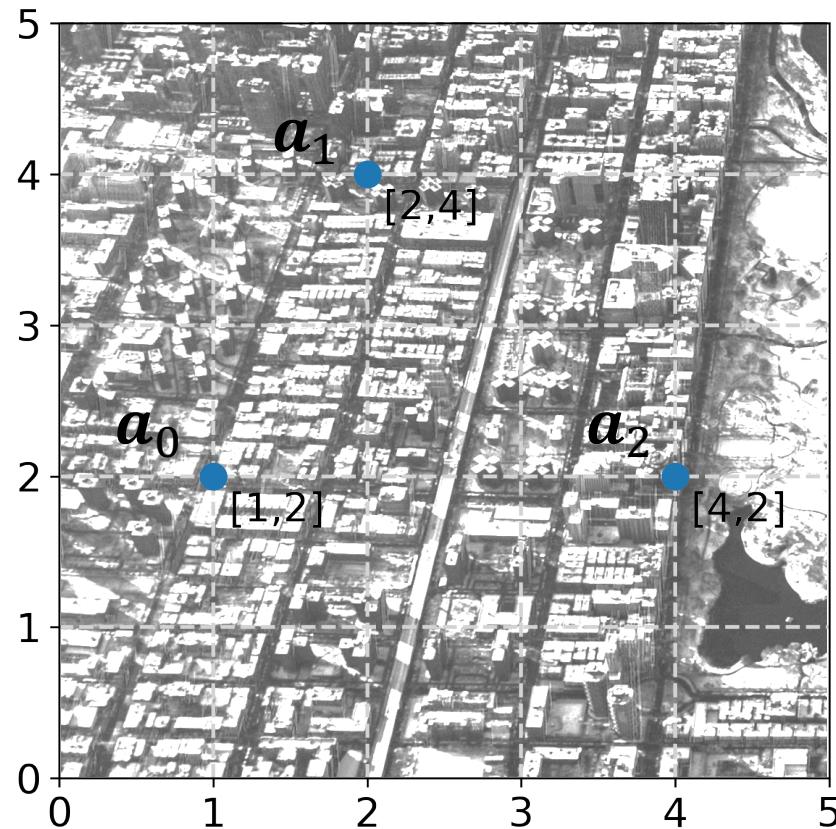


$$\begin{aligned}d_{euclidean}(a_0, a_1) &= \|a_0 - a_1\|_2^2 \\&= \sqrt{(a_{0,0} - a_{1,0})^2 + (a_{0,1} - a_{1,1})^2} \\&= \sqrt{(1 - 2)^2 + (2 - 4)^2} \\&\approx 2.24\end{aligned}$$

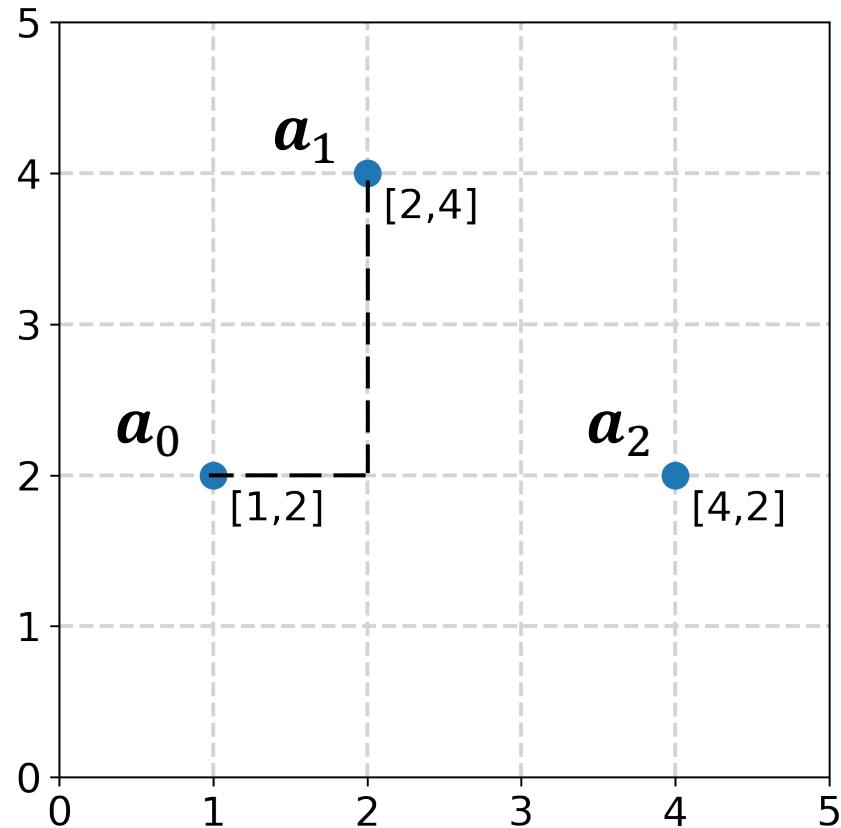
How far away is each point from a_0 ?



How far away is each point from a_0 ?



Manhattan Distance



Manhattan distance

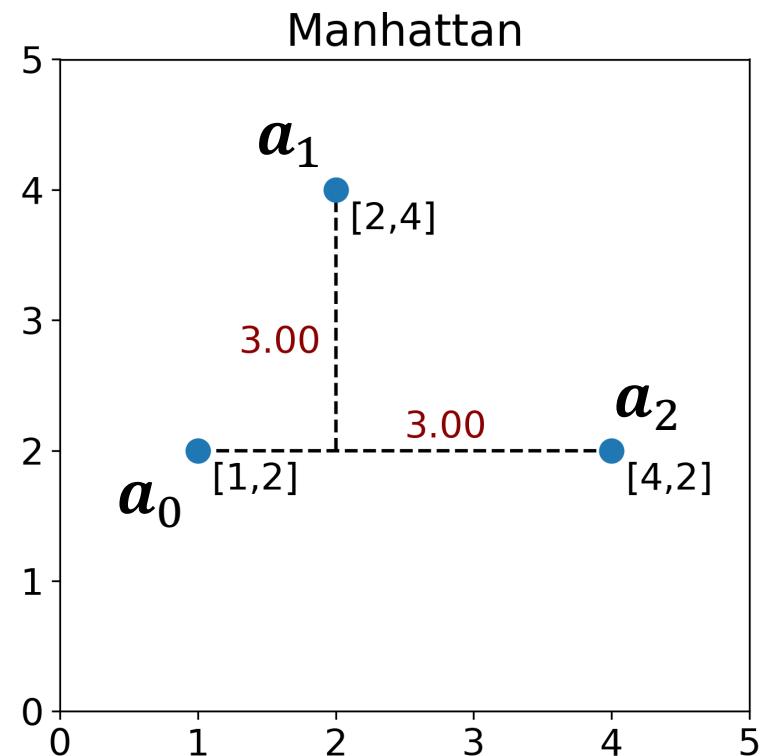
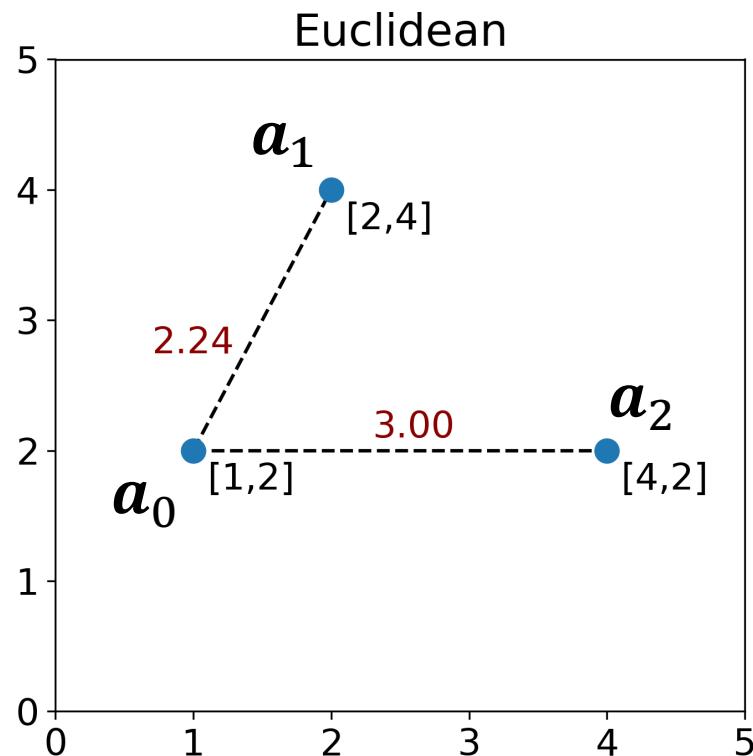
$$d_{manhattan}(a_0, a_1) = |a_{0,1} - a_{1,1}|$$

$$= |a_{0,0} - a_{1,0}| + |a_{0,1} - a_{1,1}|$$

$$= |1 - 2| + |2 - 4|$$

$$= 3$$

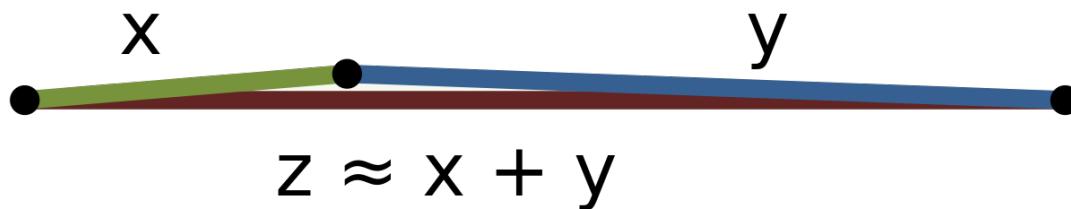
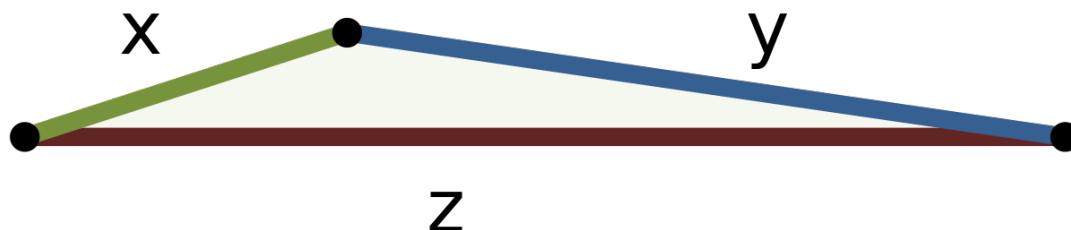
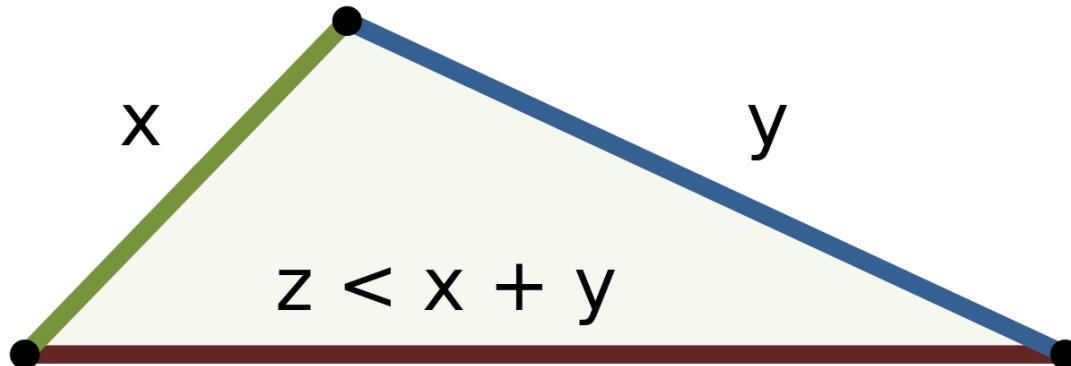
How far away is each point from a_0 ?



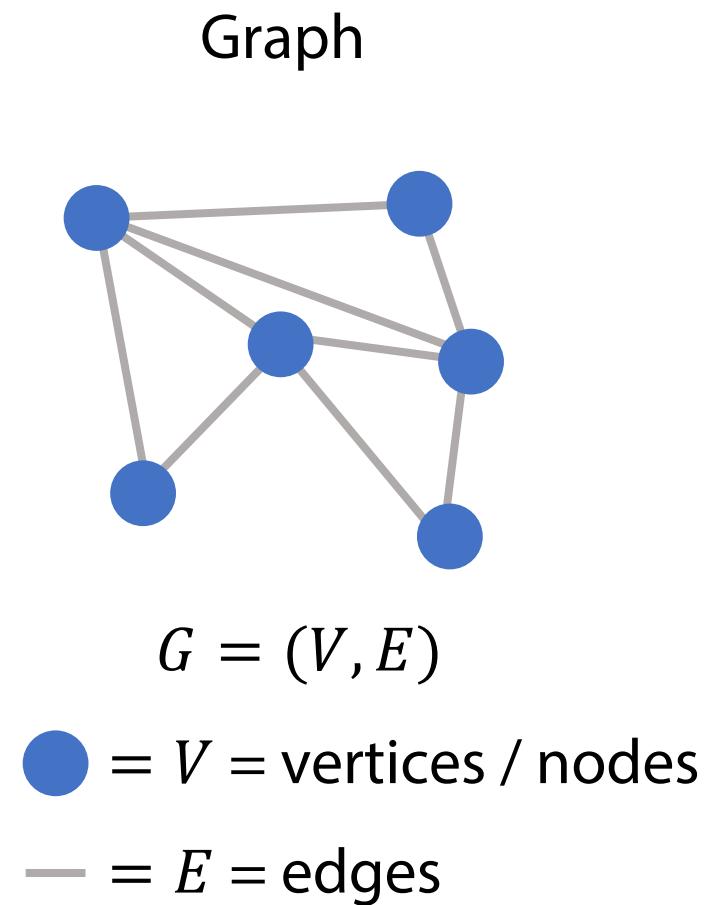
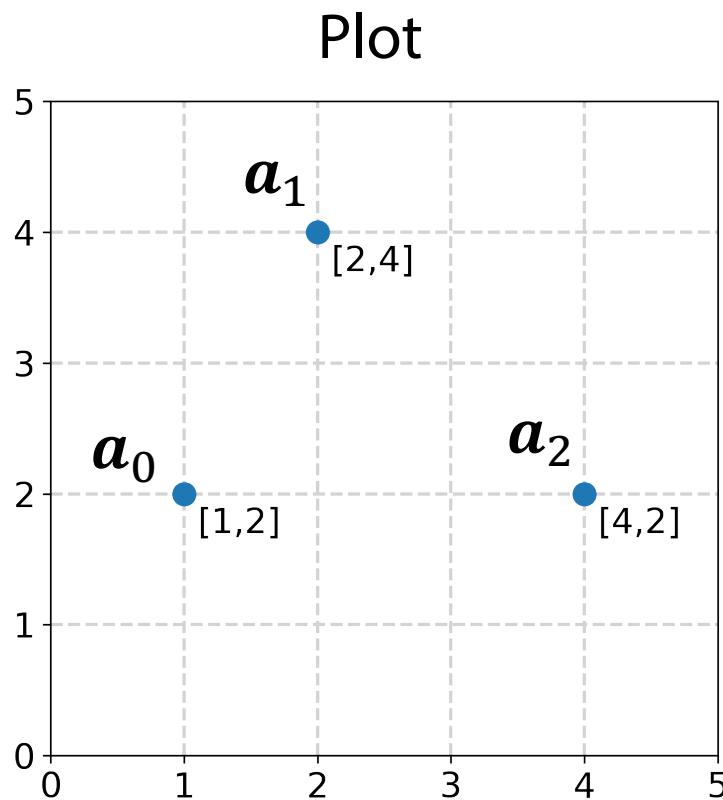
Distances

- There are many ways to measure distance
 - Hamming, Euclidean, Cosine
- Distances are functions that take two points and return a real number that is **positive or 0**
- Distances can be any function that is:
 - **Symmetric:** $\text{dist}(a \rightarrow b) = \text{dist}(b \rightarrow a)$
 - **Non-negative:** $\text{dist}(a \rightarrow b) \geq 0$
 - **Follow triangle inequality:** $\text{dist}(a \rightarrow c) \leq \text{dist}(a \rightarrow b) + \text{dist}(b \rightarrow c)$

Triangle Inequality



What is a “graph”?



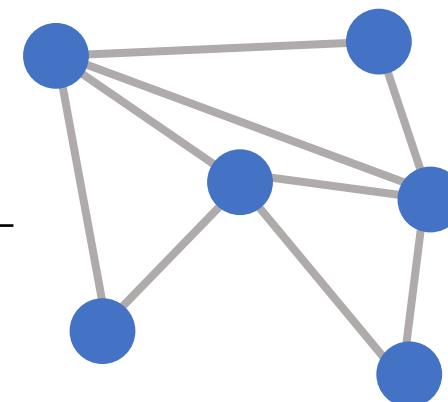


Provide an example of real-world data that could be modeled as a graph. What would be the vertices and edges?

Top

Example

Thing	Vertices	Edges
Internet	Computers	Network connections



$$G = (V, E) \quad \bullet = V = \text{vertices} \quad - = E = \text{edges}$$

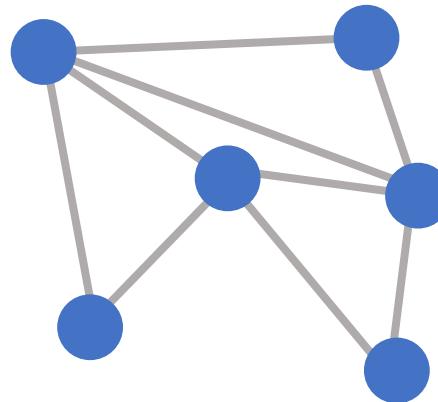


What is a “graph”?

Things that can be modelled as a graph

Thing	Vertices	Edges
Internet	Computers	Network connections
Traffic	Intersections	Roads
Social network	People	Friendships
Cell similarities	Cells	Similarity relationships

Graph

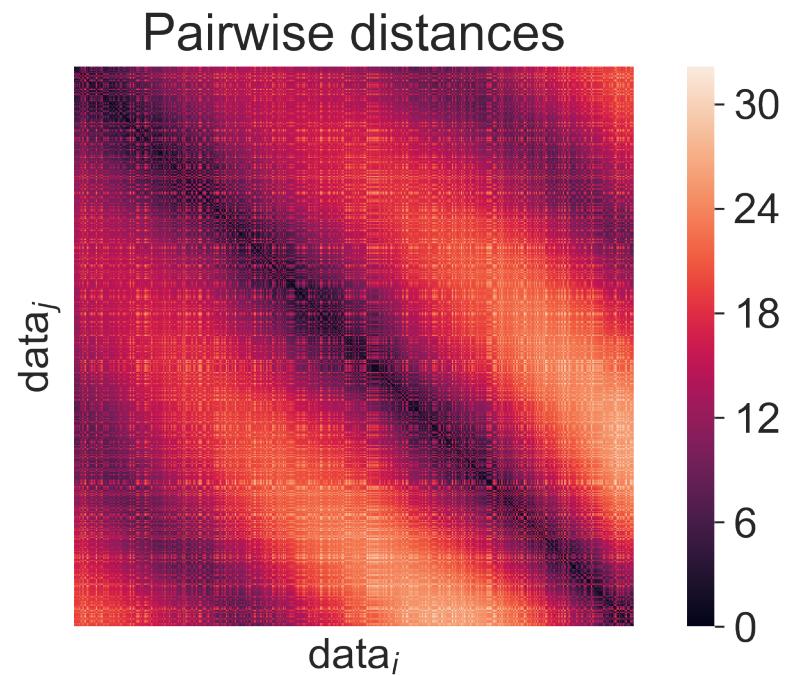
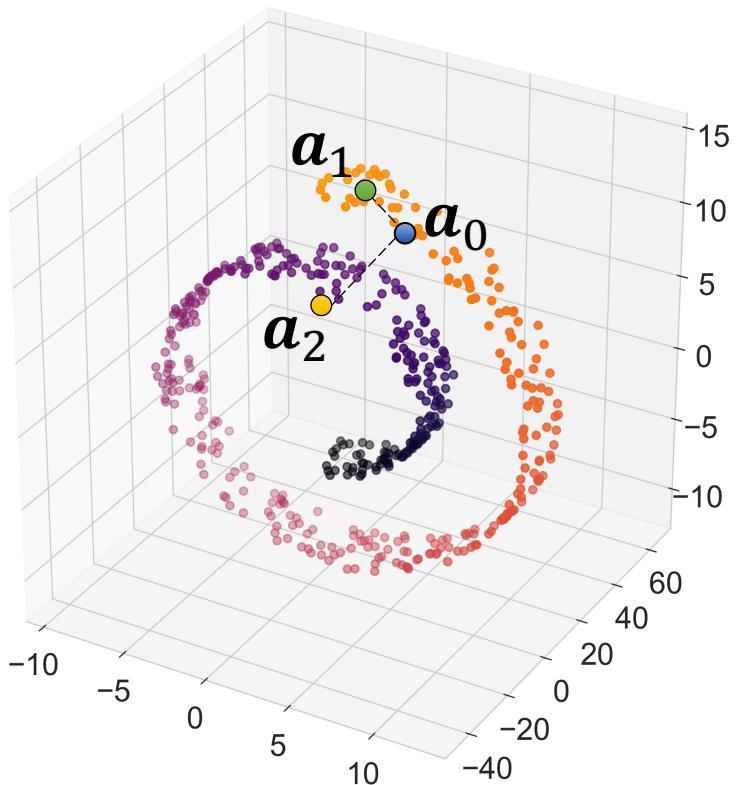


$$G = (V, E)$$

● = V = vertices

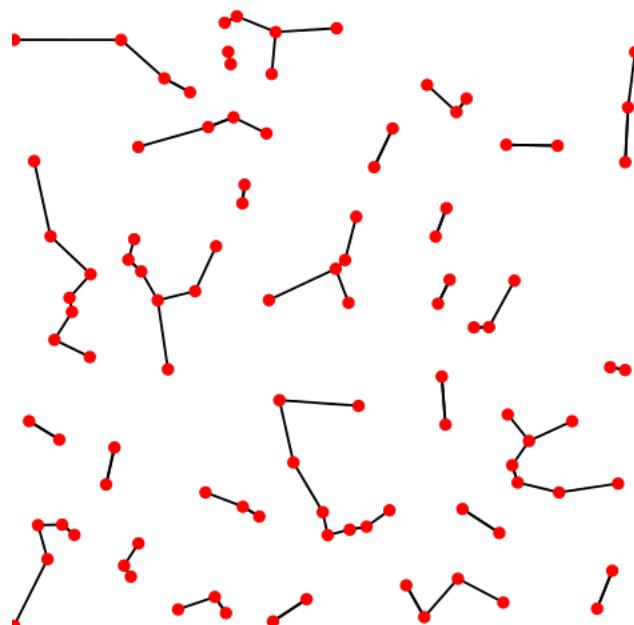
— = E = edges

How do we quantify close (local) versus far (global) in a dataset?



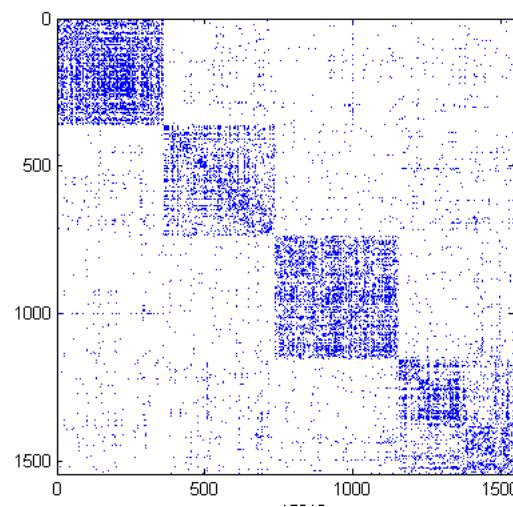
K Nearest Neighbors Graph

- To turn data into a graph, connect the K closest neighbors of a graph
- Measure closeness with your favorite distance

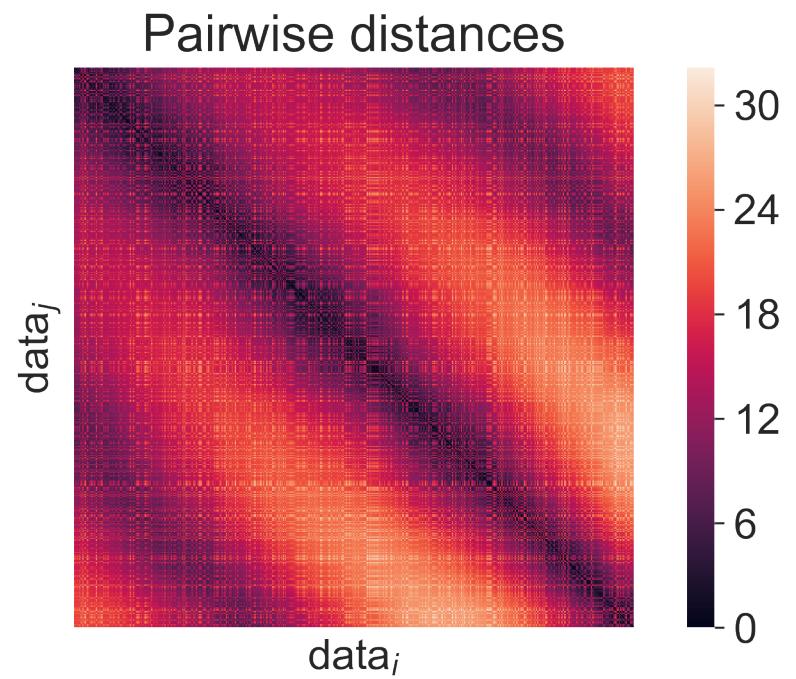
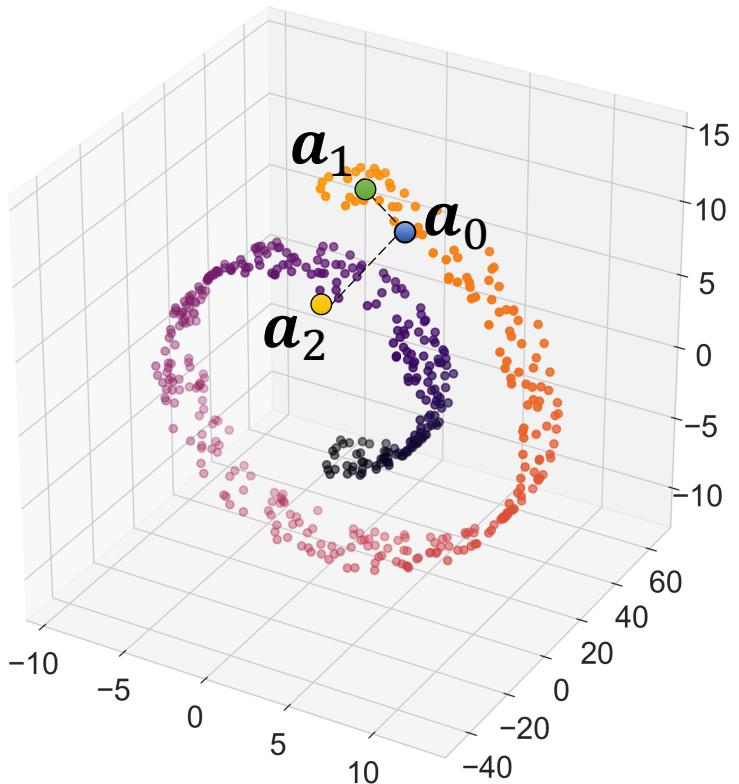


Affinity Matrix as Graph

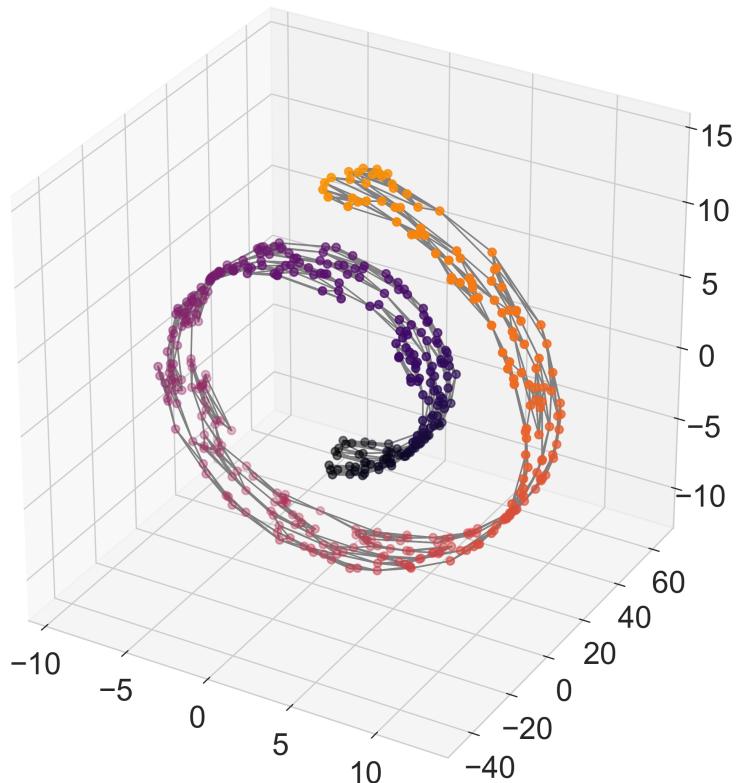
- Fully connected graph with edges weighted by an affinity (negatively proportional to distance)



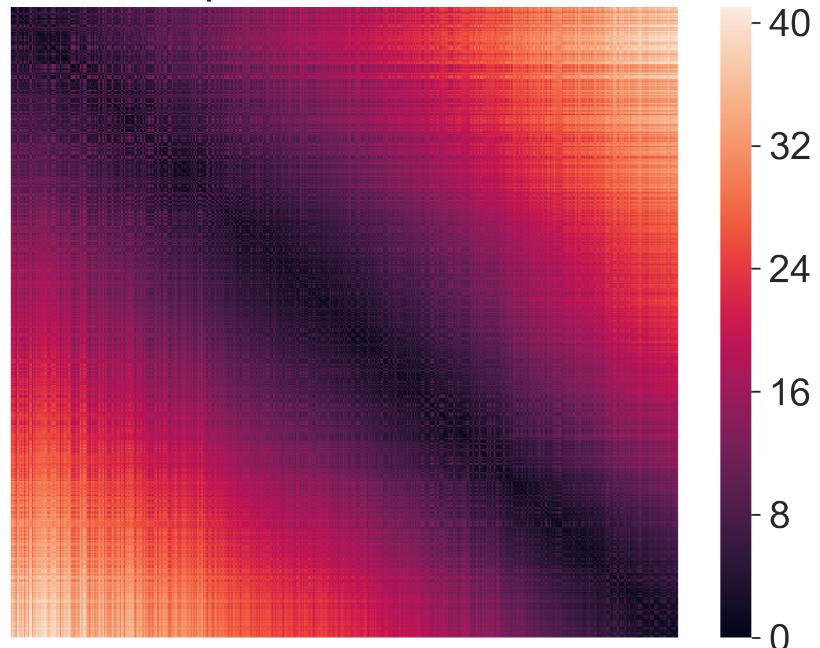
Swiss Roll



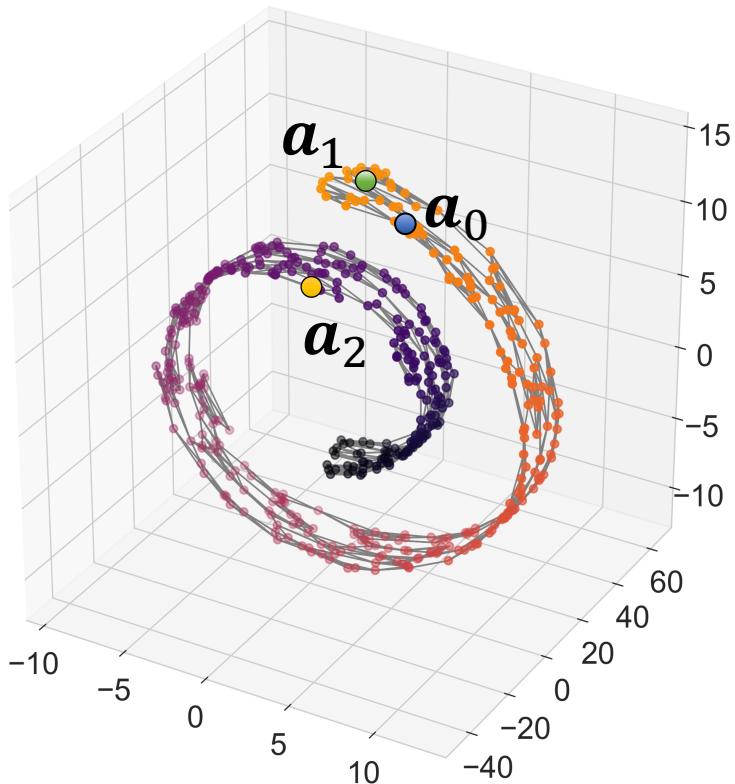
Swiss Roll K-NN



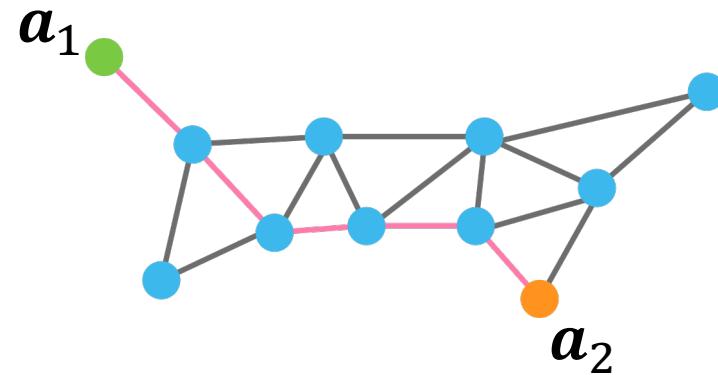
k-NN Graph Geodesic Distances



Paths or “walks” on graphs reveal dominant data manifold directions

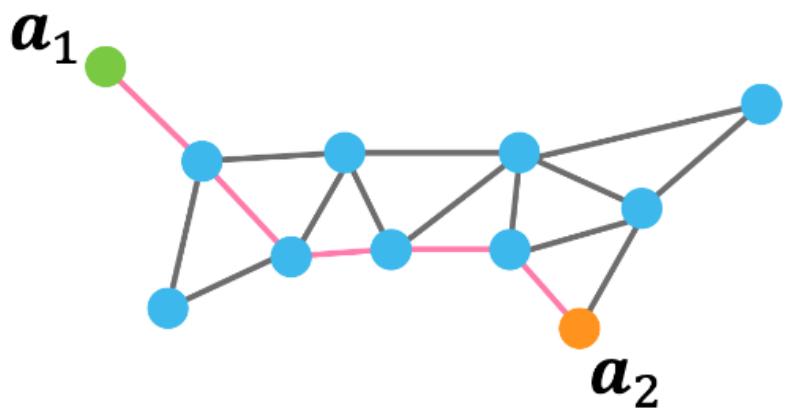


Shortest path between points



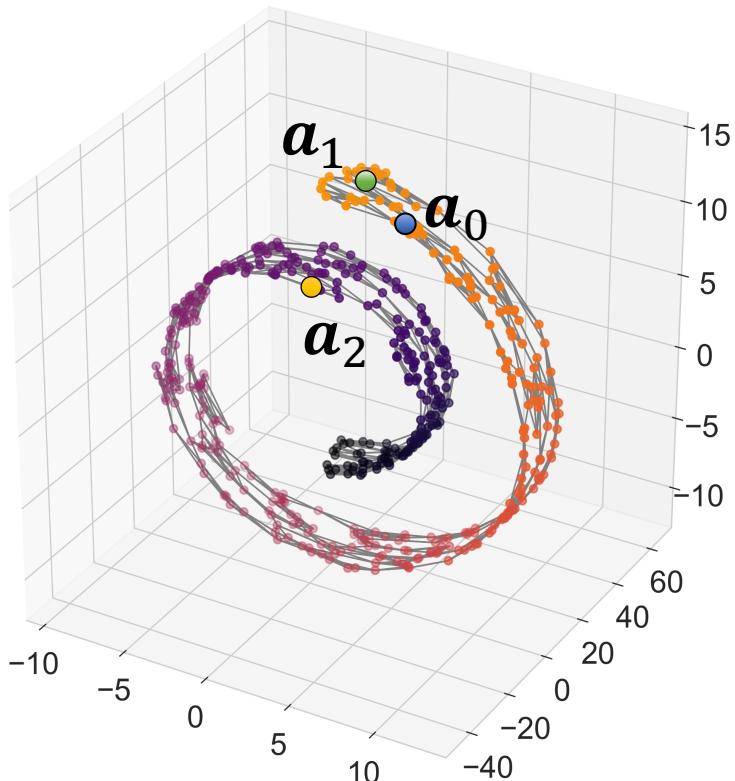
What is the shortest path distance between a1 and a2?

Shortest path between points

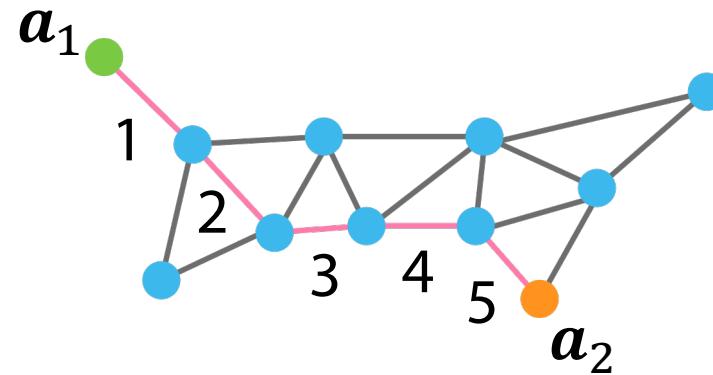


4
5
6
7
8

Graph walks approximate manifold distances

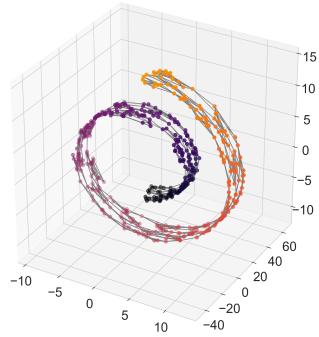


Shortest path between points

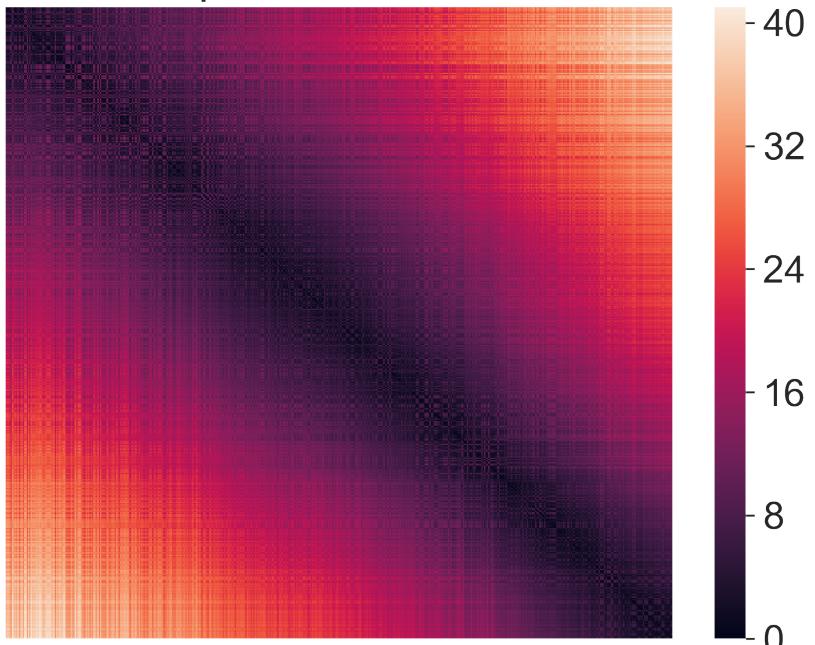


$$\begin{aligned} d_{geodesic}(a_1, a_2) &= \text{shortestpath}(a_1, a_2) \\ &= 5 \end{aligned}$$

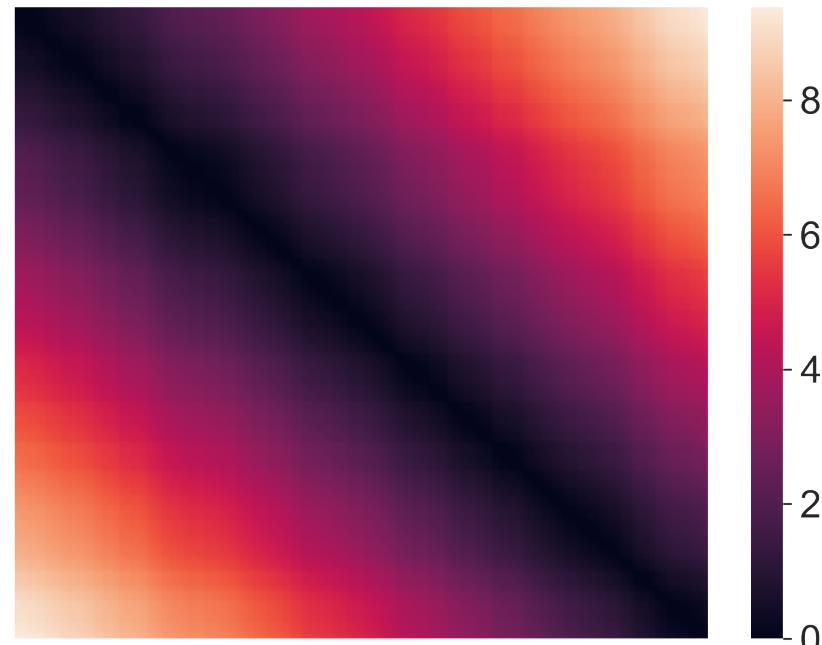
Graph walks approximate manifold distances



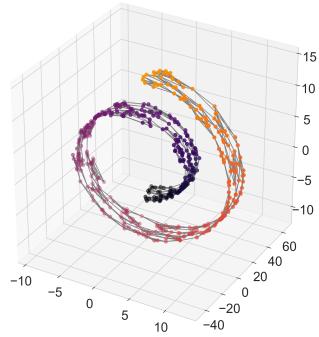
k-NN Graph Geodesic Distances



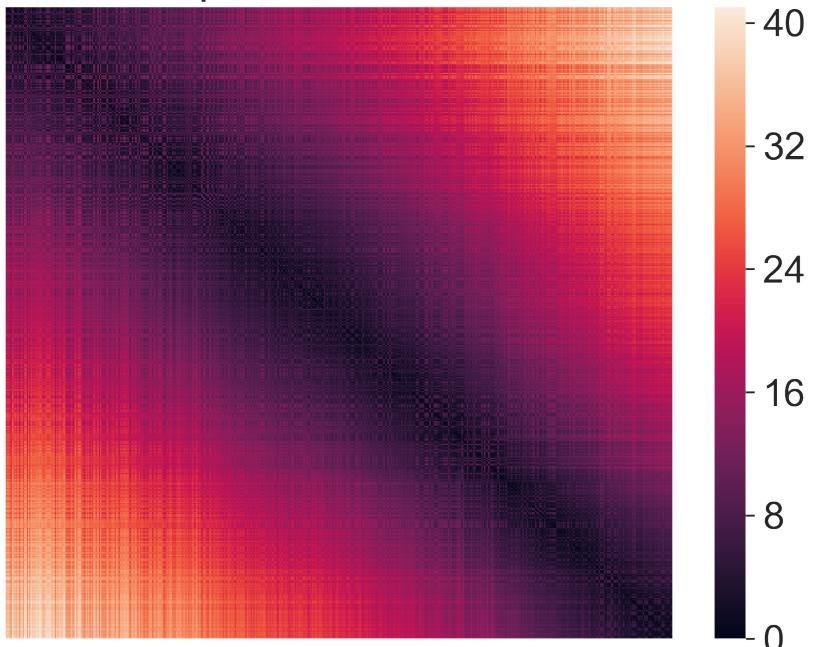
Manifold Distances



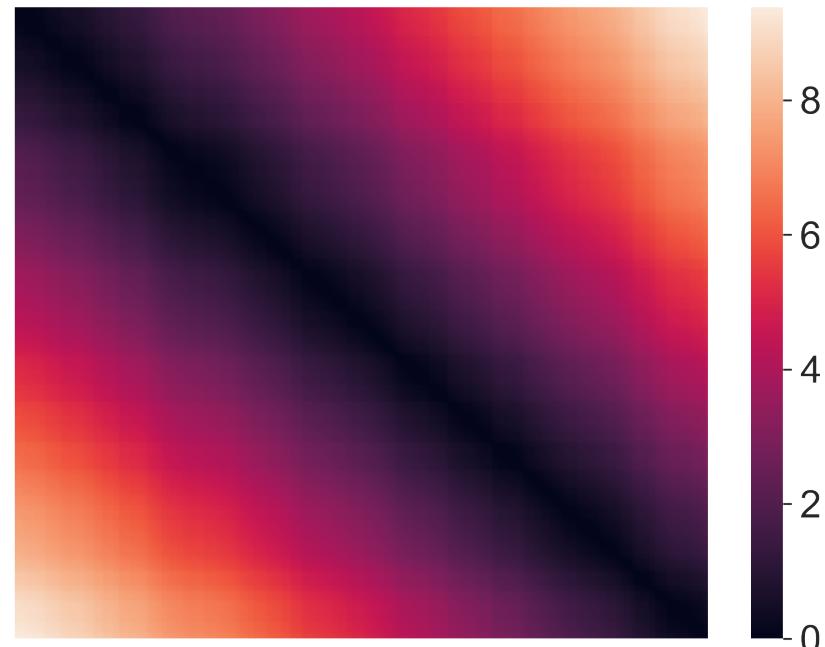
Graph walks approximate manifold distances



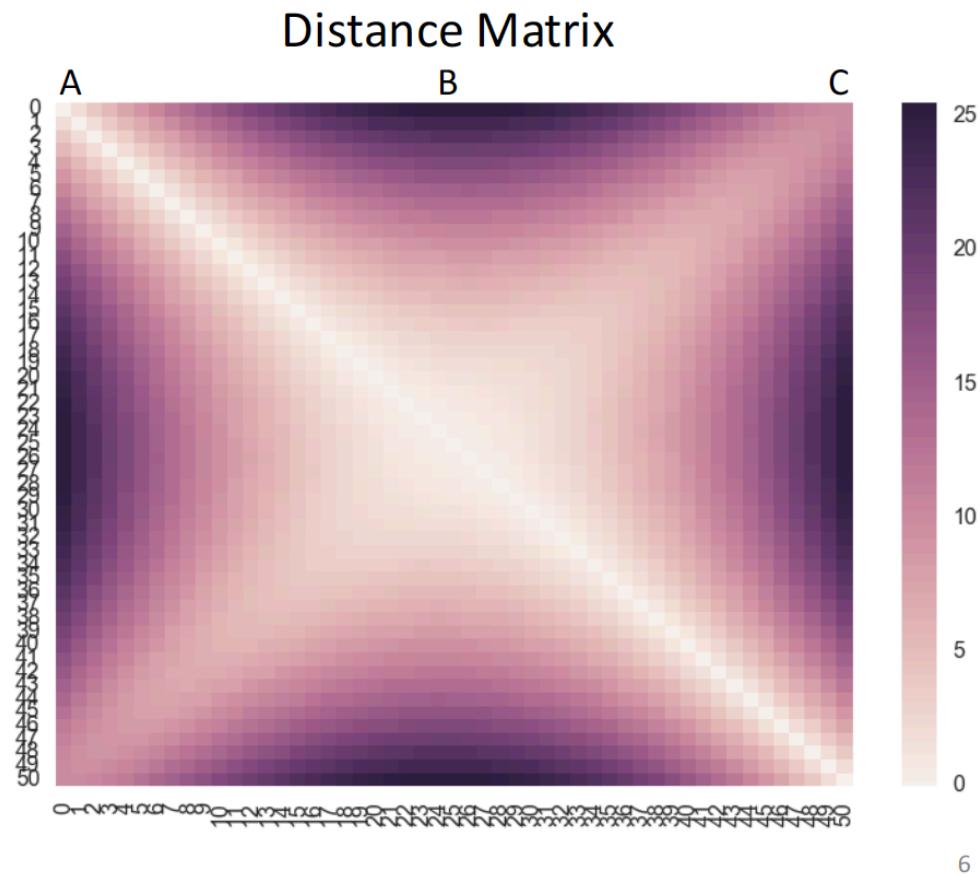
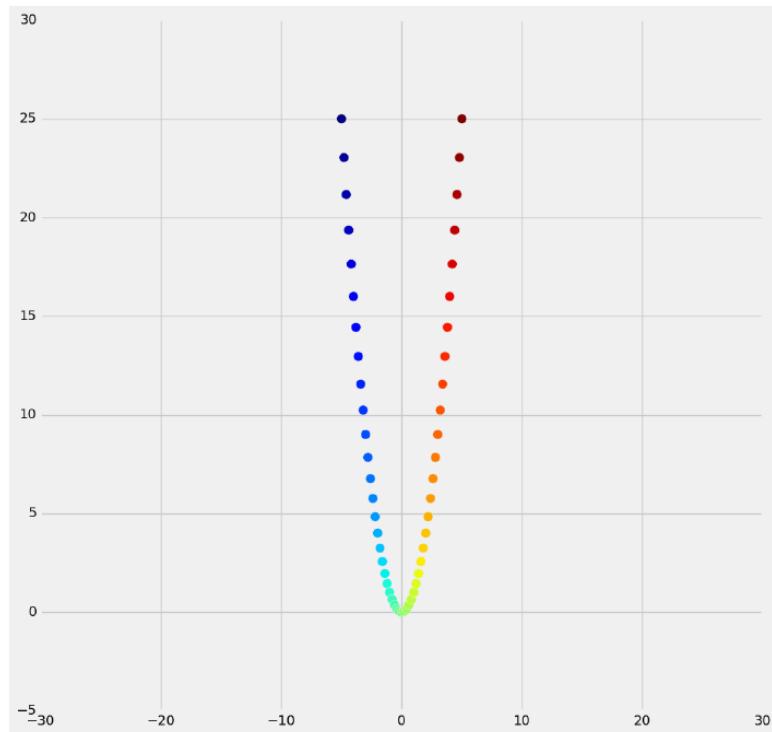
k-NN Graph Geodesic Distances



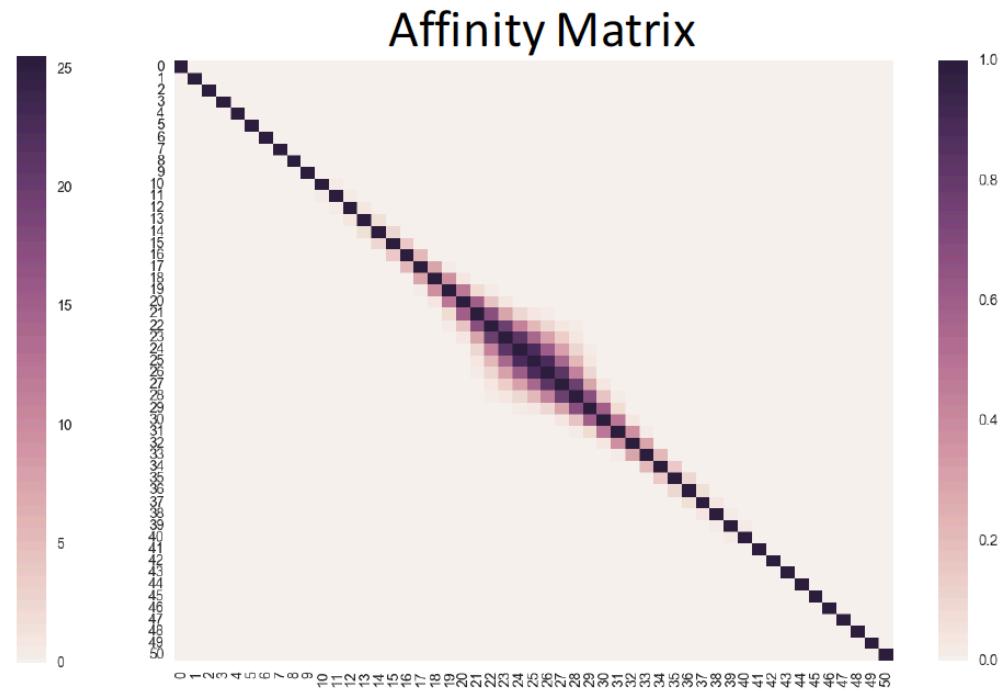
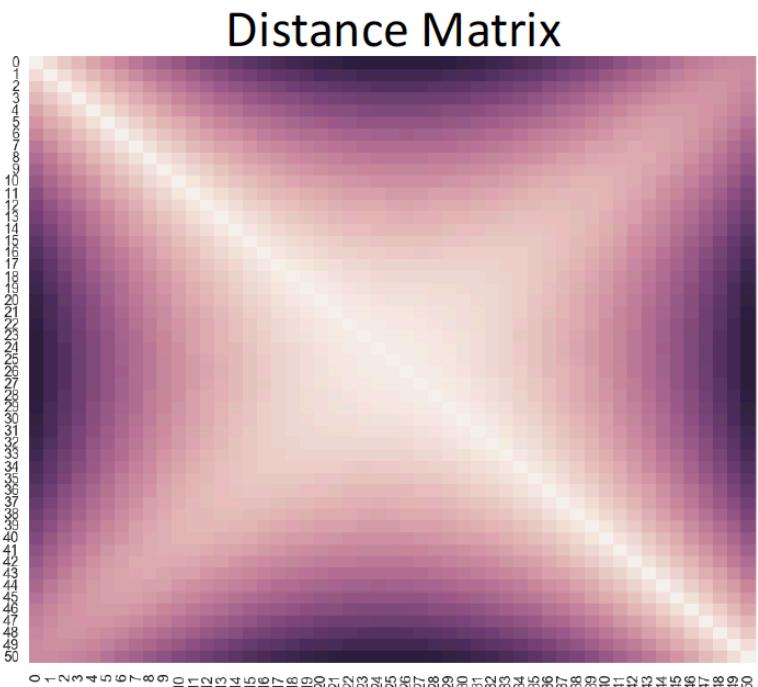
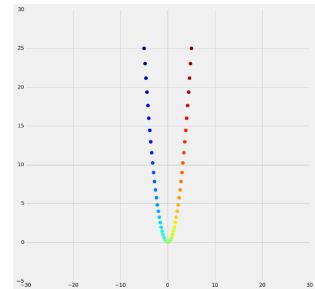
Manifold Distances



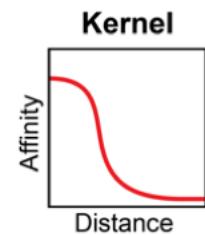
**Kernel functions measure
“similarity” or “affinity” on a graph**



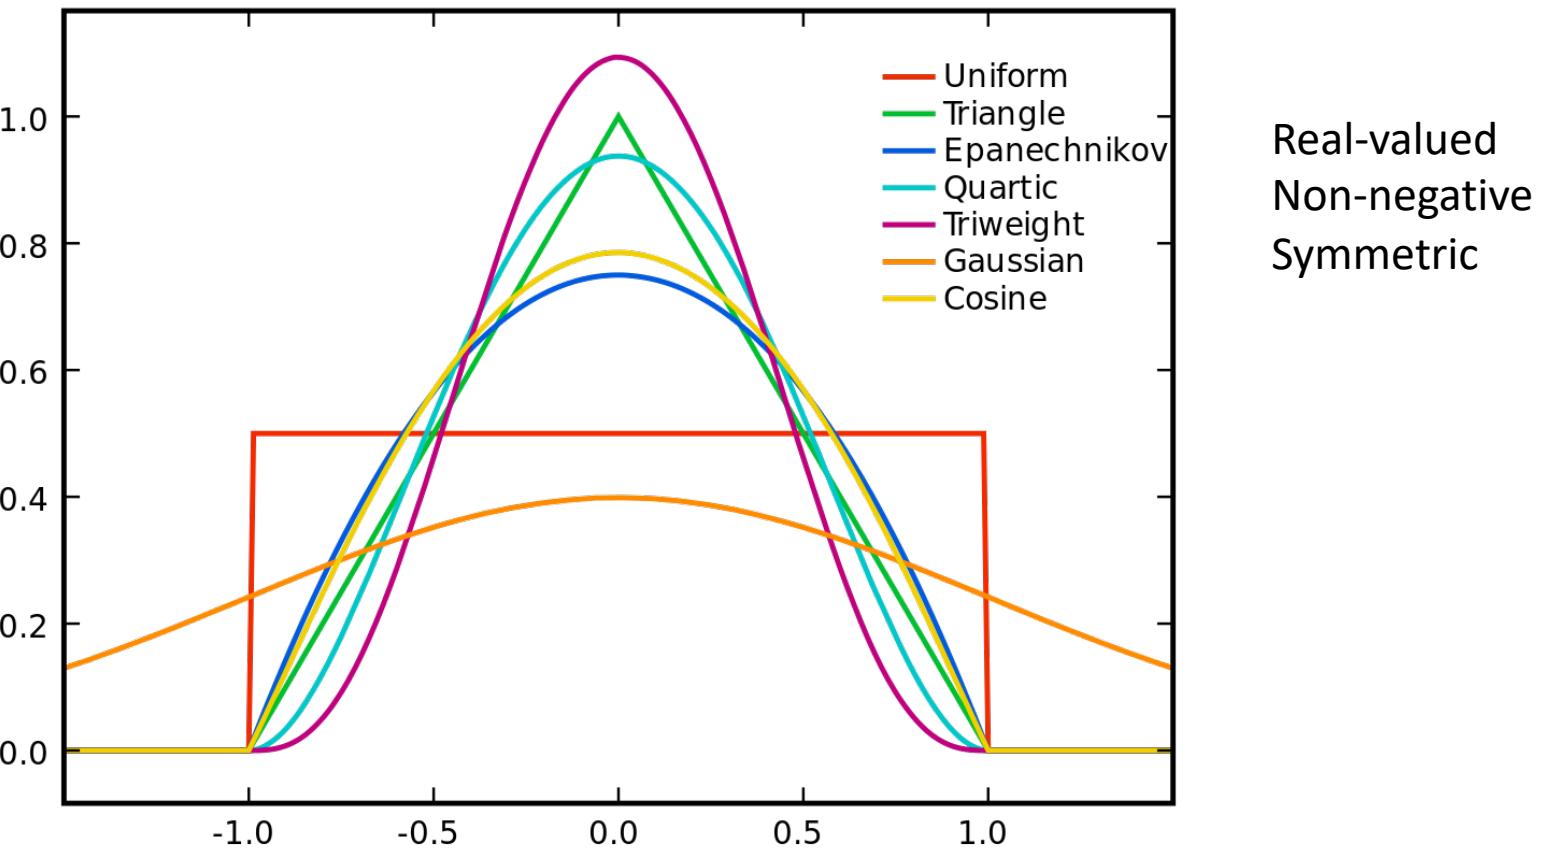
Affinity is the inverse of distance + locality



$$Affinity_{i,j} = s_{i,j} = \exp\left(-\frac{dist(x_i, x_j)^2}{2\sigma^2}\right)$$

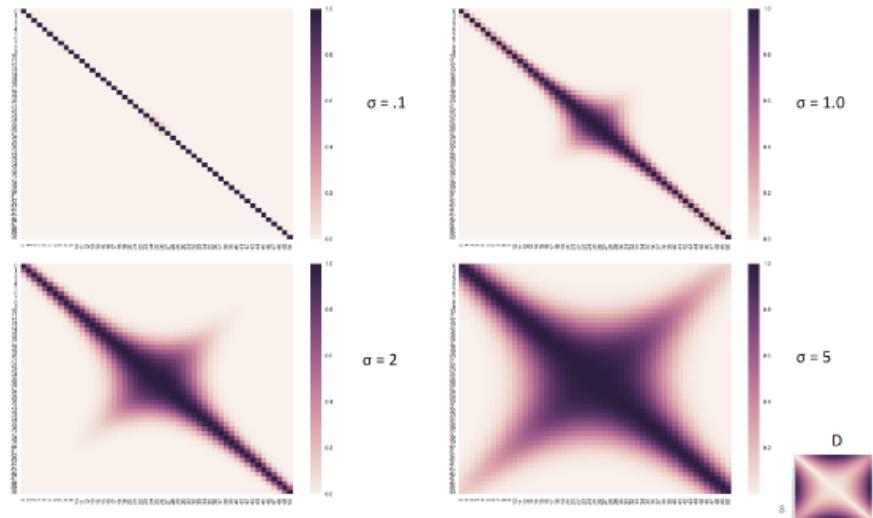


Distance to Affinity via Kernels



Affinities correlations in this hidden hypothetical space

Which value of sigma produces the graph with the least global connectivity



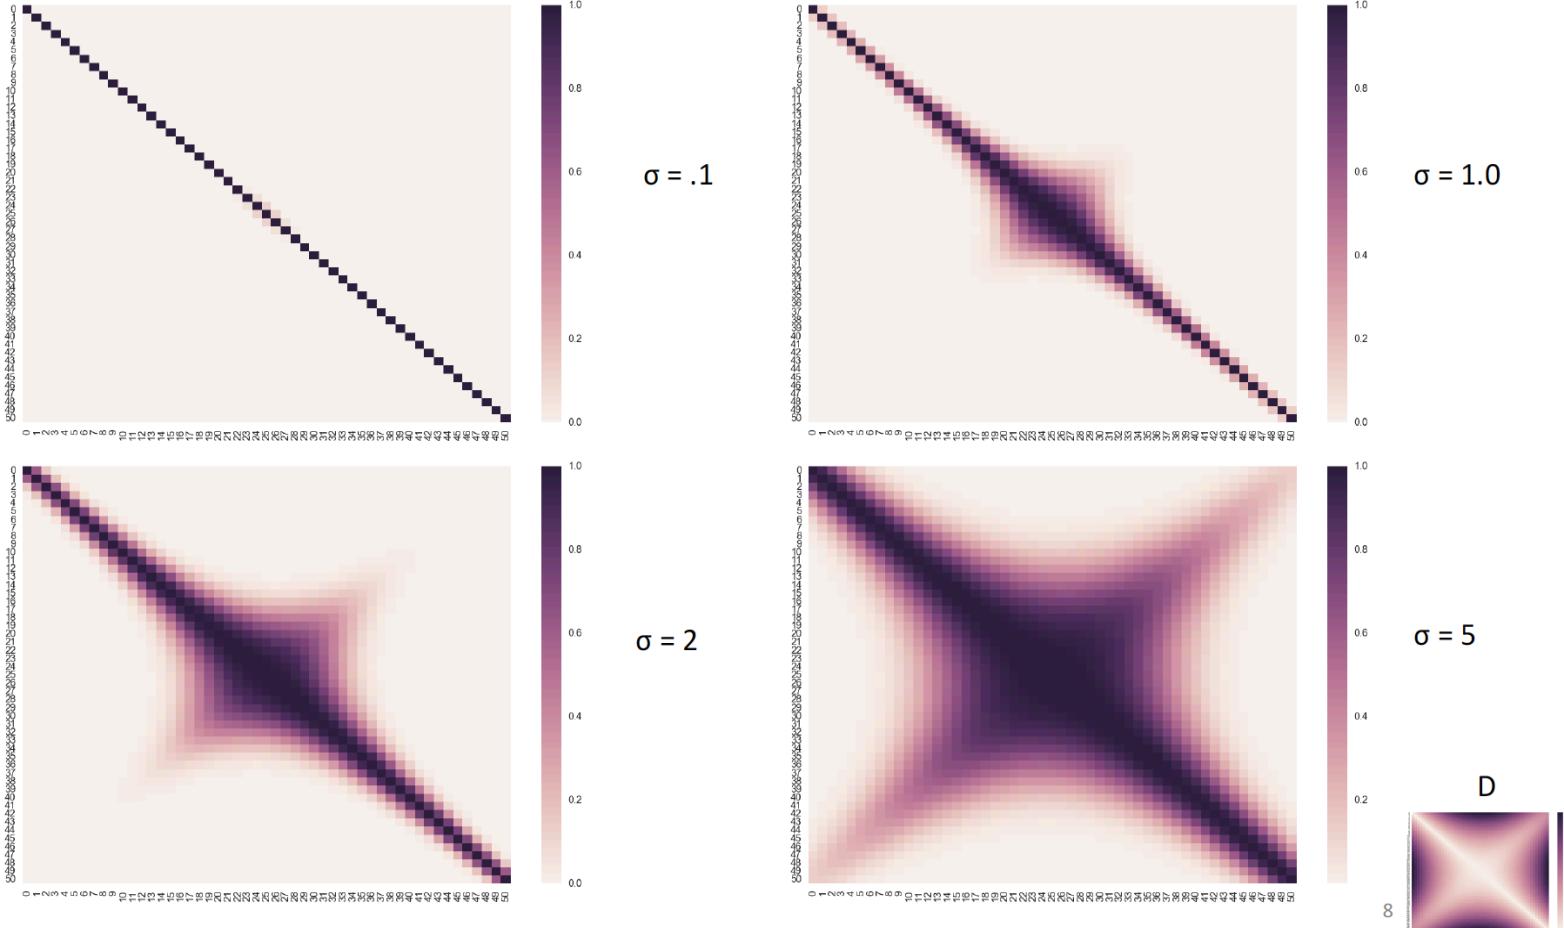
$\sigma = 0.1$

$\sigma = 1$

$\sigma = 2$

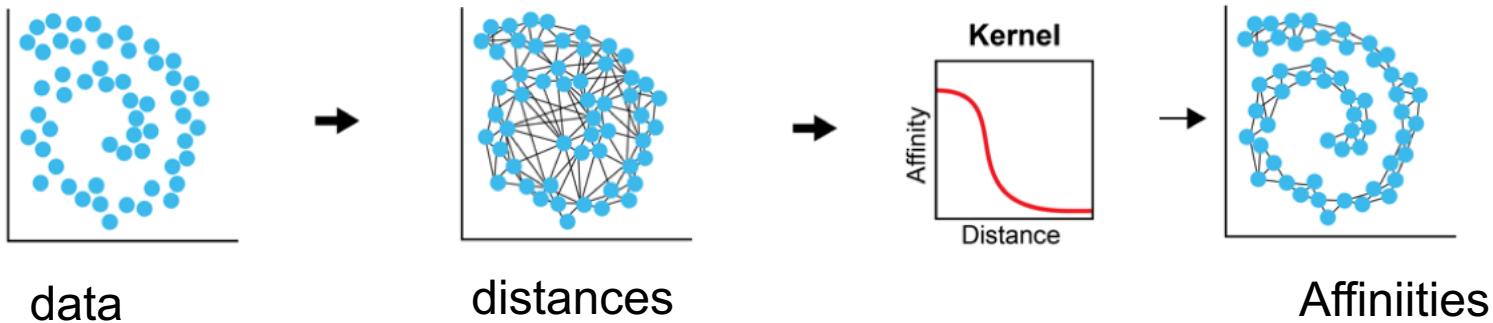
$\sigma = 5$

Effect of Changing Kernel Width



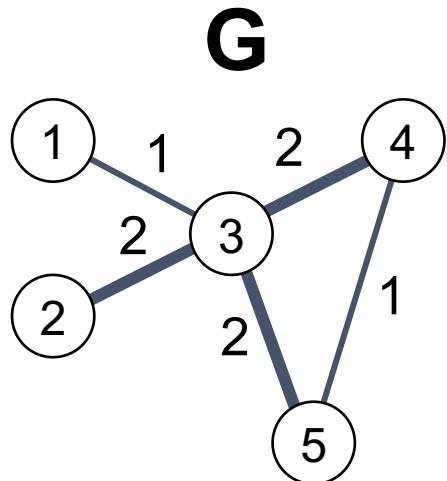
$$s_{i,j} = \exp\left(-\frac{\text{dist}(x_i, x_j)^2}{2\sigma^2}\right)$$

Kernel PCA



- Use eigenvectors of an **affinity matrix** instead of covariance matrix
- The family of methods called kernel PCA:
- Specific variants include:
 - Laplacian Eigenmaps
 - Diffusion Maps

The graph Laplacian is a matrix representation of a graph



A adjacency matrix

$$\mathbf{A} = \begin{vmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 \\ 1 & 2 & 0 & 2 & 2 \\ 0 & 0 & 2 & 0 & 1 \\ 0 & 0 & 2 & 1 & 0 \end{vmatrix}$$

D degree matrix

$$\mathbf{D} = \begin{vmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 7 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 3 \end{vmatrix}$$

Laplacian matrix

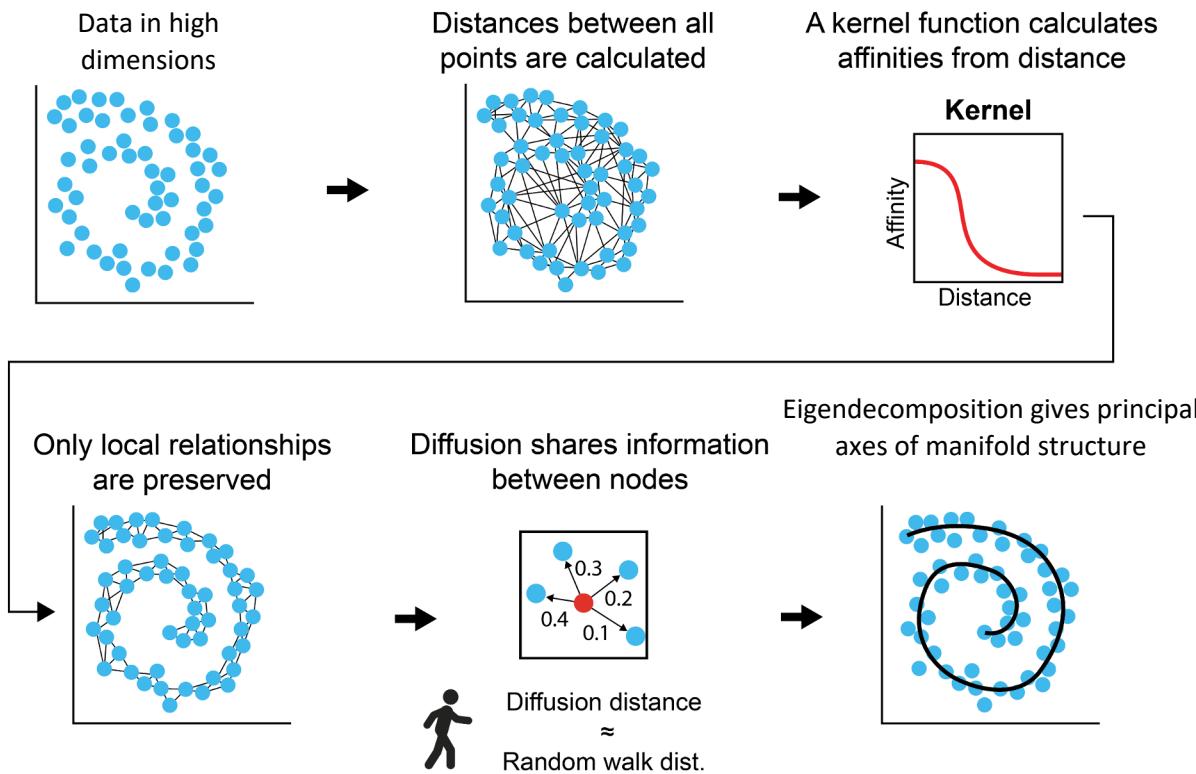
$$\mathcal{L} = \mathbf{D} - \mathbf{A} = \begin{vmatrix} 1 & 0 & -1 & 0 & 0 \\ 0 & 2 & -2 & 0 & 0 \\ -1 & -2 & 7 & -2 & -2 \\ 0 & 0 & -2 & 3 & -1 \\ 0 & 0 & -2 & -1 & 3 \end{vmatrix}$$

Laplacian Eigenmaps

$$\mathcal{L} = \mathbf{D} - \mathbf{A} = \begin{vmatrix} 1 & 0 & -1 & 0 & 0 \\ 0 & 2 & -2 & 0 & 0 \\ -1 & -2 & 7 & -2 & -2 \\ 0 & 0 & -2 & 3 & -1 \\ 0 & 0 & -2 & -1 & 3 \end{vmatrix}$$

Eigenvalues of Laplacian matrix form Laplacian Eigenmaps

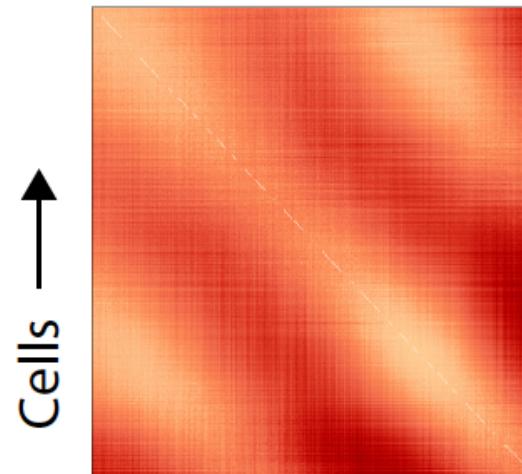
Diffusion Maps



Distance Matrix



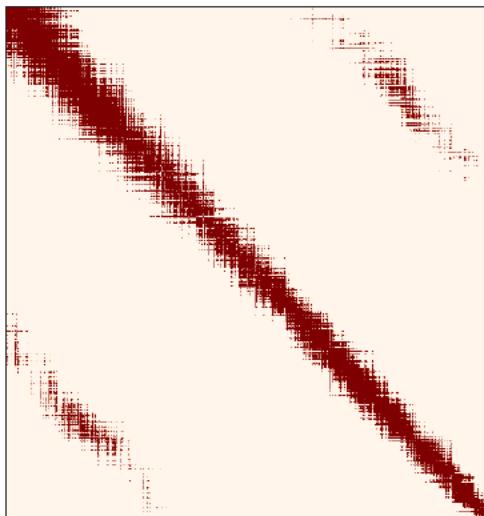
Distance Matrix



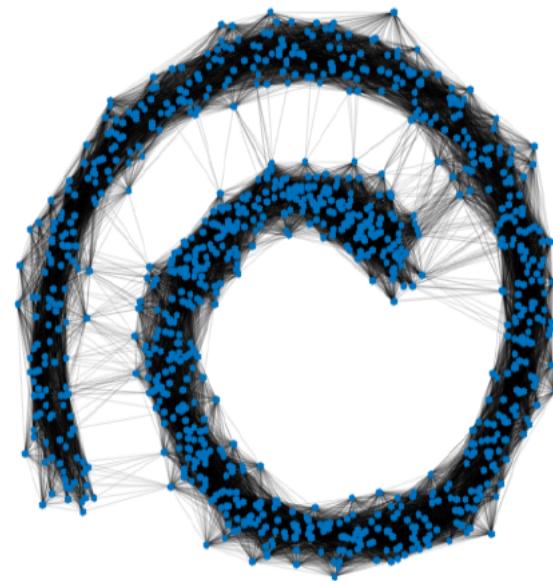
Entries are distances

$$D(x_i, x_j) = \sqrt{||x_i - x_j||}$$

Affinity Matrix



(Gaussian Kernel)



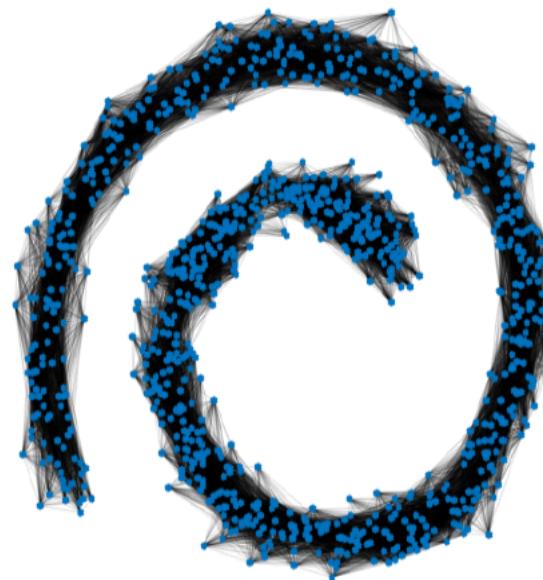
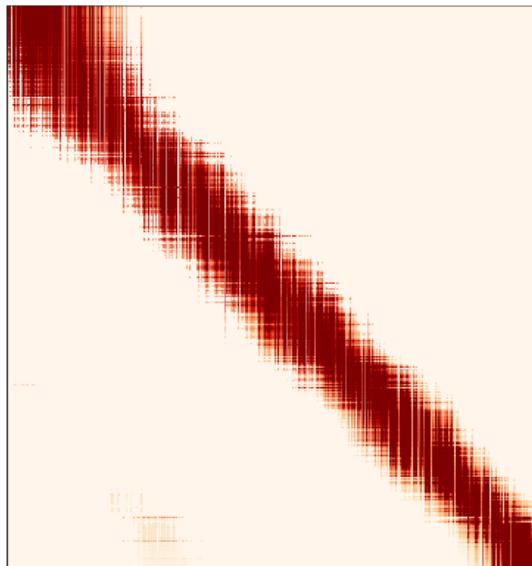
Graph Representation

Entries are affinities

$$A(x_i, x_j) = \exp\left(-\frac{D(x_i, x_j)^2}{\sigma}\right)$$

Powered Markov Matrix

Powered Markov Matrix

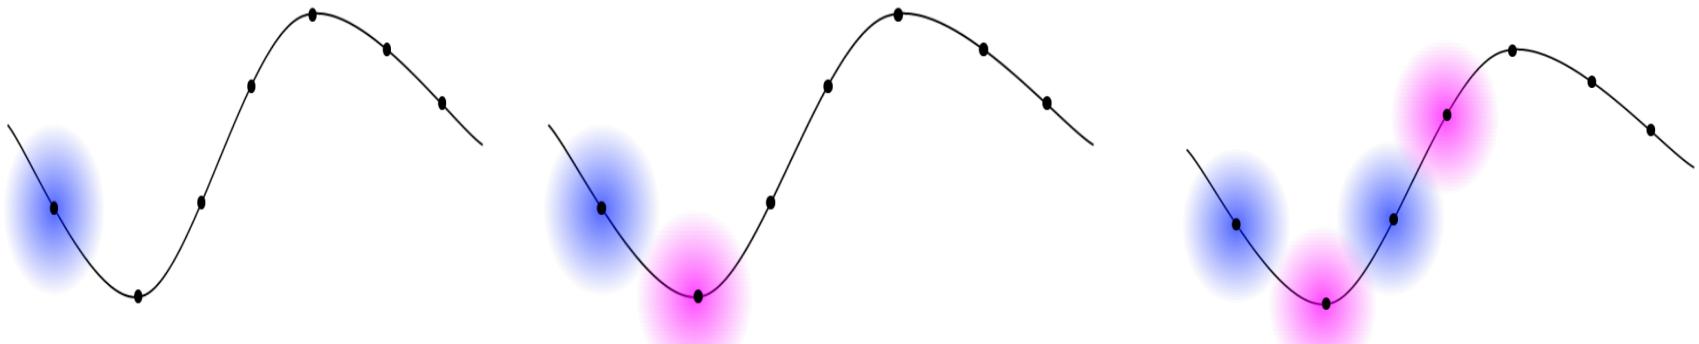


Entries are row normalized affinities

$$M(x_i, x_j) = \frac{A(x_i, x_j)}{\sum_j A(x_i, x_j)}$$

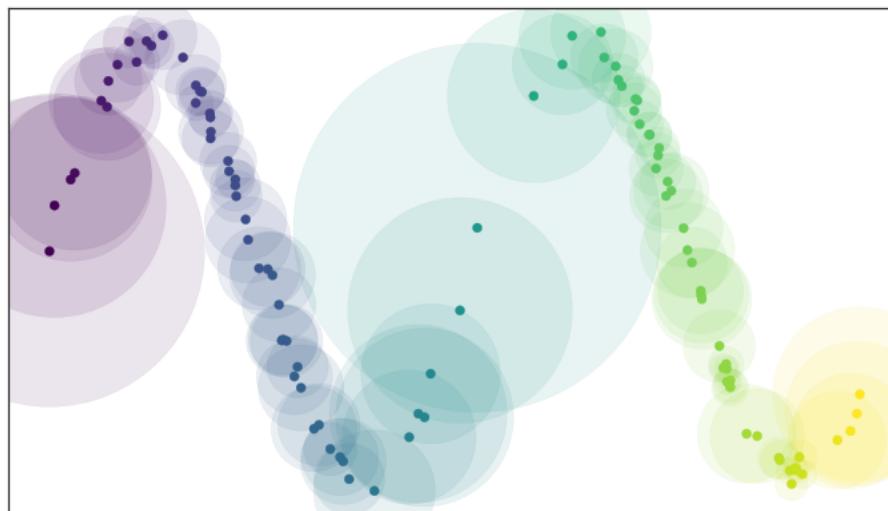
What does Powering Do?

- Smooths and infers connectivity through data



Adaptive neighborhoods

- Neighborhoods can adapt to the density of the local area
- This actually helps diffusion discover manifold geometry by avoiding “dense ditches”



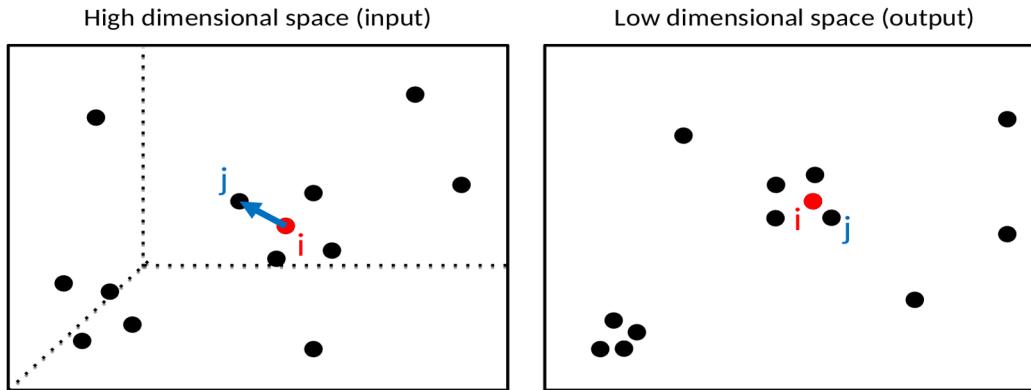
Diffusion maps pros and cons

- ✓ Find intrinsic dimensions within noisy data
 - ✓ Denoises data well
- ✓ Captures trajectories or clusters in the data
- ✓ Distances are meaningful in diffusion maps
- ✓ Robust in that runs are the same
 - ✓ Deterministic method
- ✗ Separates each component into different dimensions
 - ✗ Each eigenvector is roughly one trajectory or one cluster ---bad for a 2D visualization (but good for clustering!)
- ✗ Slow runtime---due to eigendecomposition of a large matrix

2D Dimensionality Reductions

- These methods preserve information from an affinity matrix in 2-3 dimensions exactly
- There are as many eigenvectors as datapoints, so taking only the top 2 misses information
 - **tSNE/UMAP:** greedy method for preserving near neighbors in 2D
 - **PHATE:** creates a new distance matrix from a distance between diffusion probabilities of points and squeezes that variation into 2D

tSNE (and UMAP)



Goal: Maintain the local neighborhood of each cell

Steps:

1. Place points randomly in a 2D plane
2. Move points to locally optimal location via gradient descent

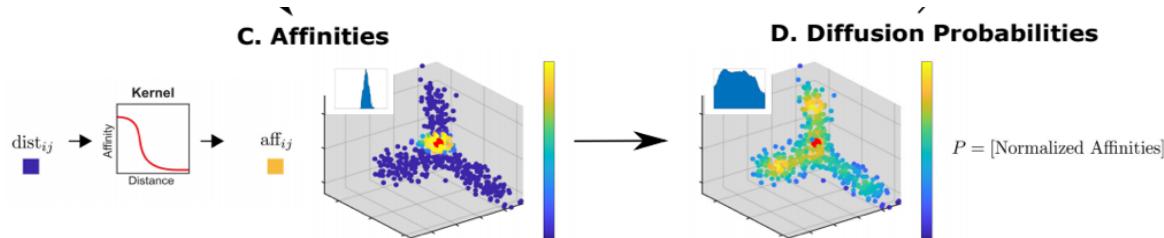
$$\frac{\delta C}{\delta y_i} = 2 \sum_j (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j})(y_i - y_j).$$

tSNE/UMAP pros and cons

- ✓ Captures information in 2 D
- ✓ Preserves local neighborhoods
- ✓ Shows number of points in a cluster by flattening out density (i.e. does not “overplot”)
- ✓ UMAP fast due to SGD and lack of normalization
 - ✓ tSNE can be slow for large number of points and low dimensions
- ✗ Does not preserve any type of density or distance
- ✗ Does not capture global structure of data
- ✗ Does not denoise data
- ✗ Based on local minima, solution can change run-to-run

PHATE:

- Goal: to preserve manifold structure
- Step 1. Each cell is represented as a t-step diffusion probability to other cells



- Step 2: Re-represent each point by its diffusion probabilities to all other points

$$C_1 = [p_{11} \ p_{12} \ \dots \ p_{1n}]$$

$$C_2 = [p_{21} \ p_{22} \ \dots \ p_{2n}]$$

...

$$C_n = [p_{n1} \ p_{n2} \ \dots \ p_{nn}]$$

- Step 3. Define a new distance between points with this new representation

$$C_1 = [p_{11} \ p_{12} \ \dots \ p_{1n}]$$

.....

$$C_i = [p_{i1} \ p_{i2} \ \dots \ p_{in}]$$

...

$$C_j = [p_{j1} \ p_{j2} \ \dots \ p_{jn}]$$

...

$$C_n = [p_{n1} \ p_{n2} \ \dots \ p_{nn}]$$

$$\text{dist}_{ij} = \sqrt{\|\log P_i - \log P_j\|^2}$$

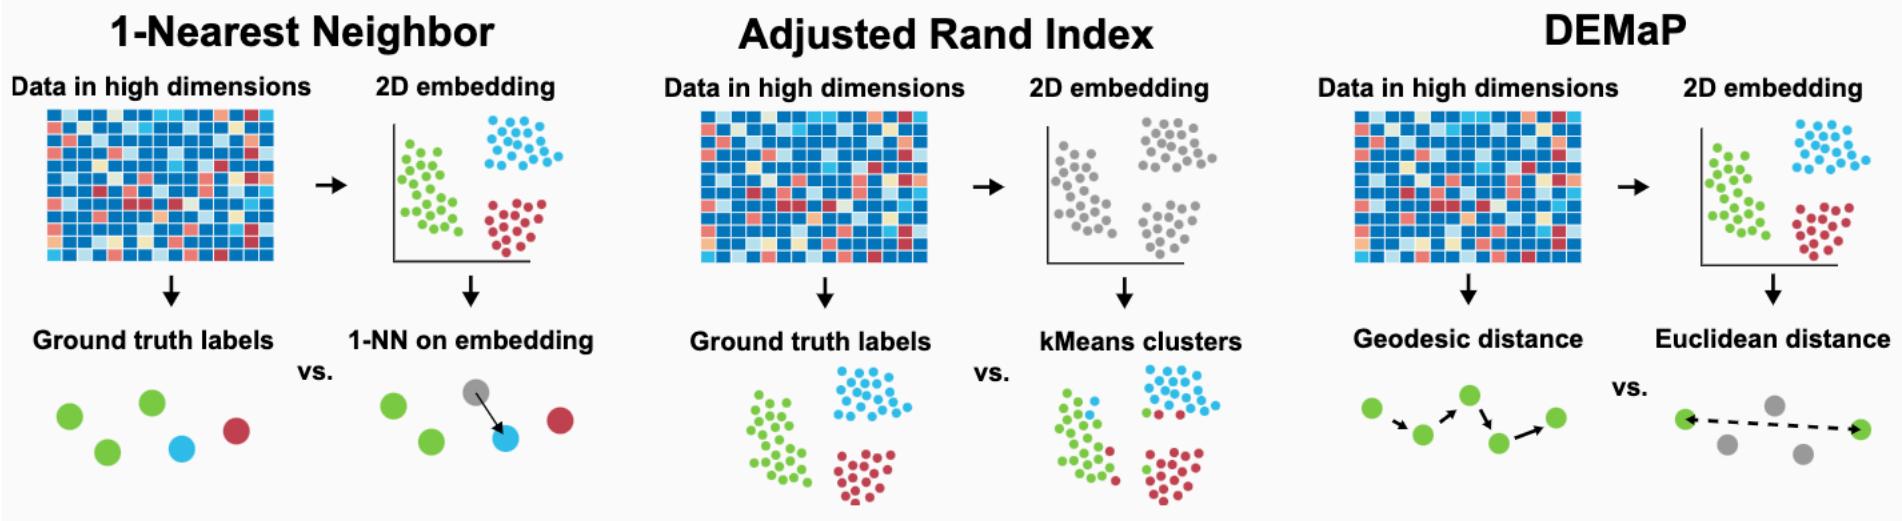


- Step 4: Embed with multidimensional scaling
- This reorganizes information in a diffusion map into 2D

PHATE pros and cons

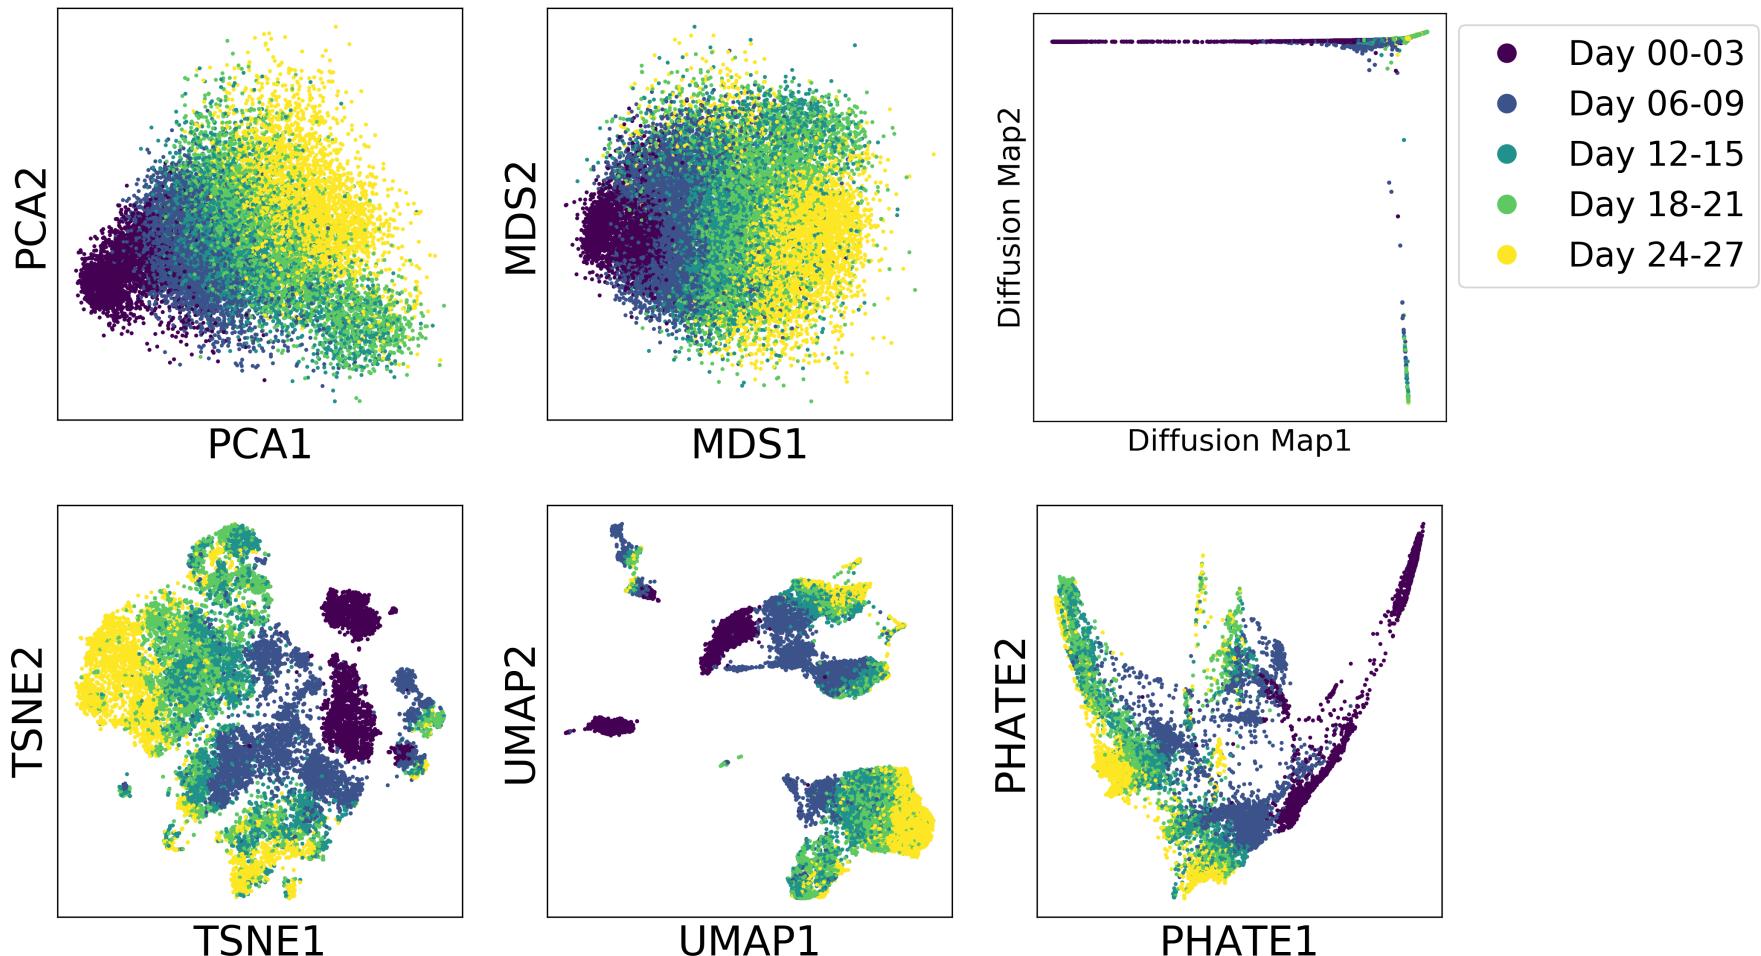
- ✓ Captures information in 2 D
- ✓ Denoises data for clean visualization
- ✓ Preserves local and global manifold structure
- ✓ Fast due to landmarked diffusion tricks
- ✓ Distances are meaningful in PHATE (preserves Denoised Manifold Affinity)
- ✓ Robust, same solution run-to-run
- ✗ Overplots data similar points can be very close
- ✗ Shows “horseshoe” shapes due to MDS

Quantitative Metrics

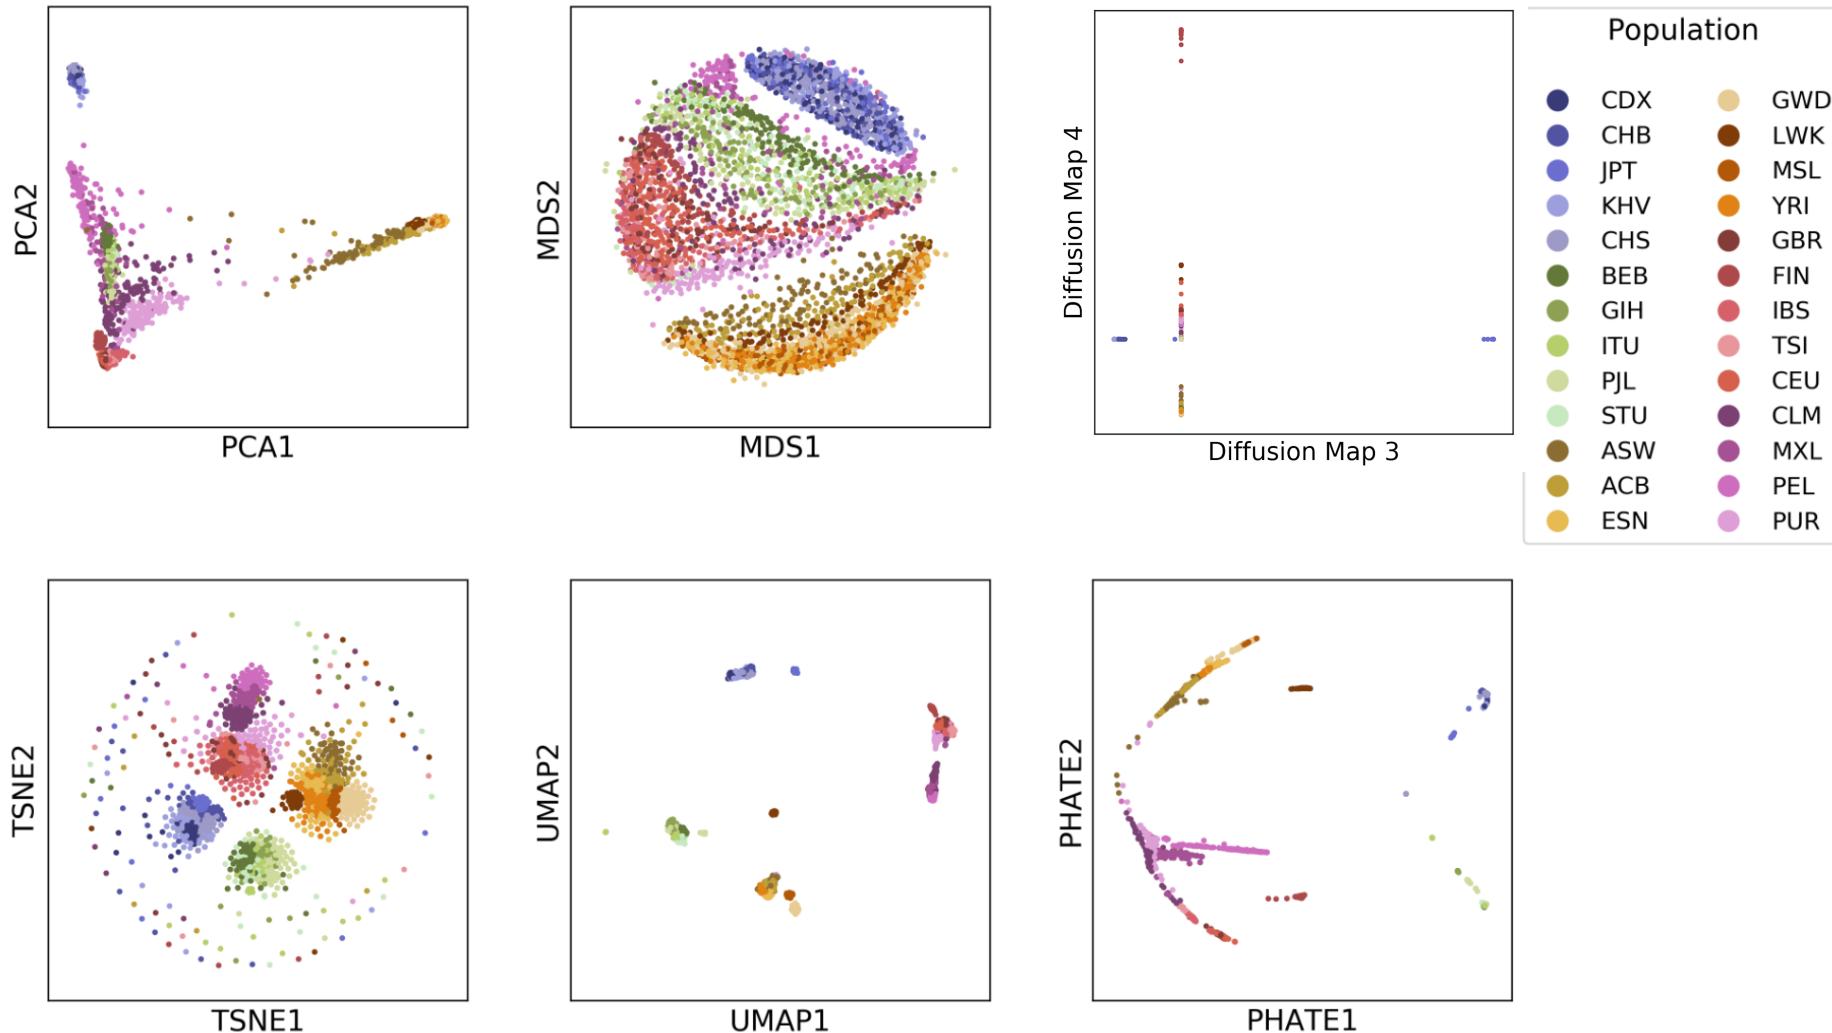


PHATE tends to do better on DeMAP, and ARI tSNE/UMAP on 1-NN, sometimes on ARI on certain Datasets

Visual Information Extraction: Embryoid Body



1000 Genomes



When poll is active, respond at **PollEv.com/yaleml**

Text **YALEML** to **22333** once to join

What visualization would be good for understanding the heterogeneity in T cells measured after an infection?

When poll is active, respond at **PollEv.com/yaleml**

Text **YALEML** to **22333** once to join

What visualization might be good for deciding if a dataset has a single dominant trajectory of progression?

Conclusions

- Dimensionality reduction and visualization helps to identify structure in data
- PCA uses eigenvectors to identify linear combinations of features that explain maximum variance
- Manifold learning identifies non-linear patterns and structure
- Graphs can be used to approximate manifolds and are helpful for modelling single cell data
- Laplacian Eigenmaps and diffusion maps visualize data using dominant non-linear directions in data manifold
- PHATE summarizes denoised manifold information in 2D
- UMAP, tSNE are non-linear methods focused on preserving local neighborhoods in 2D

What questions do you have about today's lecture material?

Top